Lecture Notes
Distributed System

Hinnerk van Bruinehsen
Tobias Höppner
Tobias Famulla

SoSe 2013

Contents

Chapter 1

Verteilte Systeme/Distributed Systems

1.1 Orga

VL Di 10-12 (nicht am 23.04.)
Ue Do 10-12

1.1.1 Elektisches

- (kvv)
- Website AG
- Sakai

1.1.2 Übungen

- ca. 5 Übungsblätter, 14-tägig
- Vorträge in Gruppen über „verteilte Systeme"

1.1.3 Material/Inhalt

1. Hälfte Distributed Systems (Tanenbaum, van Steen)
    - Architektur
    - Prozesse
    - Kommunikation
    - Namen
    - Synchronisation
    - Konsistenz
    - Replikation
    - Fehlertoleranz
2. Hälfte Distributed Algorithms (Nancy Lynch)
    - synchronous network algorithms
    - network models (leader election, shortest path, distributed consensus, byzantine agreement)

- asynchronous network algorithms (shared memory, mutual exclusion, resource allocation, consensus)
- timing
- network resource allocation
- failure detectors

Chapter 2

Distributed Systems

Def: A distributed System is a collection of independent computers that appears to it's users as a single coherent system.

Characteristics:

- autonomous components
- appears as single system
- communication is hidden
- organisation is hidden
  (could be high-performance mainframe or sensor net)
- heterogenous system offers homogenous look/interface

Objectives:

- provide resources (printer, storage, computing)
    - share in a controlled, efficient way
    - grant access $\Rightarrow$ connect users and resources

2.1 Transparency

hide the fact that processes and resources are physically distributed.

Types of transparancy:

access hide differences in data representation and how a resource is accessed
location hide where a resource is located
migration hide that a resource may move to another location
relocation hide that a resource may be moved to another location while in use
replication hide that a resource is replicated
concurrency hide that a resource may be shared by serveral competitive users
failure hide the failure and recovery of a resource

transparancy is desireable, but not always perfectly possible
tradeoff between transparancy and complexity, maintainablility and performance

2.2 Open System

- service interfaces specified using Interface Definition Language (IDL)
- service specification as text

2.3 Scalability

is an important property
- scalable in size (number of nodes)
- scalable in geographic spread
- scalable in administration

2.4 Problems

- centralized services
- centralized data
- centralized algorithms

Scaling techiques
- use only asynchronous communication
- distribution, split components
- replication of components

pitfalls
1. reliable network
2. secure network
3. homogenous network
4. constant topologgy
5. zero latency
6. infinite bandwith
7. zero transport cost
8. one administrator!

Types of distributed systems
- computing systems
  - cluster computing

Figure 2.1: cluster computing

  - grid computing(virtual organisation, geographically distributed and heterogenous))
- distributed inforamtion systems
  - transaction processing systems (database)
    ACID (atomicity, consistency, isolated, durable)
  - enterprise systems

- Distributed pervasive systems
  small, wireless, adhoc, no administration
  home automation, health systems, sensor networks

Why do we need distributed systems?

- performance
- distribution inherent
- reliability
- incremental growth (scalability)
- sharing resources

Chapter 3

Architectures of distributed Systems

- how to split software into components
  $\Rightarrow$ Softwarearchiticture
- how to build a system out of the components
  $\Rightarrow$ Systemarchitecture

Middleware can help to create distribution transparency

Architecturestyles:

- Layered architecture
  $\Rightarrow$ network stack, messages or data flow up and down

  - control flow between layers
  - requests down
  - reply up
- Object-based architectures

  - interaction between components
  - e.g. remote procedure calls
  - can be client-server system
- data-centered architectures

  - data is key element
  - communication over data, distributed database
  - web-systems mostly data-centric
- event-based architecture

  - publish-subscribe systems
  - processes communicates threeough events
  - publisher announces events at broker
    $\Rightarrow$ loose coupling (publisher and subscriber need not to know each other), decoupled in space
    $\Rightarrow$ scalability better than client-server, parallel processing, caching

Figure 3.1: publish subsribe system

Event-based and data-based can be combined
$\Rightarrow$ shared Data space

Figure 3.2: shared data space

3.1 System architectures

1. centralized architectures
   client - server

   (i) single point of failure
   (ii) performance (server is bottleneck)

Figure 3.3: client server simple waiting situation

   (a) communication problems

   (b) server problems

can request be repeated without harm?
$\Rightarrow$ request is idempotent
(iii) aplication layering
Layers:
1.) User interface
2.) processing
3.) data level

Figure 3.4: application layer

$\Rightarrow$ a lot of waiting
$\Rightarrow$ does not scale

2. Decentralized architectures

- vertical distribution (layering)
  different logic on different machines
- horizontal distribution
  replicated client/server operating on different data
  $\Rightarrow$ overlay-underlay hides physical structure by adding logical structure

Structured P2P architectures

- most popular technique is distributed hashtables (DHT)

- randomly 128 bit or 160 bit ke for data and nodes. Two or more keys are very unlikely

- Chord system arranges items in a ring
- data item k is assigened to node with smallest identifier id $\geq$ k

ie item 1 belongs to node 1
item 2 belongs to node 2
for each item k_i succ(k)=id
returns the name of the node k is assigened to
to find data item k the function LOOKUP(k) returns the adress of succ(k) in O(log(N))(later!)

membership management
join:
create SHA1 identifier
LOOKUP(id) = succ(id)
contact succ(id) and pred(id) to join ring

leave:
node id informs succ(id) and pred(id) and assigns it's data to succ(id)

10

Content adressable network (CAN)

- d-dimensional cartesian space
- every node draws random number
- space is divided among nodes
- every data draws identifier (coodinates) which assigns a node
- join
    - select random point
    - half the square in which id falls
    - assign item to centers
- leave
    - one node takes the rectangle
        $\Rightarrow$ reassign rectangles periodically

Unstructured P2P Network

- random graph
- each node maintains a list of c neighbours
- partial view or neighbourhood list with age
- nodes exchange neighbour information
    active thread
    select peer

    PUSH
    select c/2 youngest entries+myself
    send to peer

    PULL
    receive peer buffer
    construct new partial view
    increment age

    passive thread
    recieve buffer from peer

    PULL:
    select c/2
    send to peer
    construct new partial view increment age

Chapter 4

PeerSim

Chapter 5

Processes

processes
-execution of program
-processor creates virtual processor

-for each program everyting is stored in process table

-transparent sharing of resources,(processor, memory) separation
-each virtual processor has it's own independent adress space
-process switch is expensive, (save cpu context, pointers, translation lookaside buffer (TLB), memory management unit (MMU))
-perhaps even swaps to disk, if memory exhausted

threads
-several threads share CPU
-thread context has little memory information, perhaps mutex lock
-threads avoid blocking application (e.g. spreadsheet,computation of dependent cells, intermediate backup)
-thread switch is fast

-user level threads allow parallel computation of program sections
I/O or other blocking system calls block all threads, but thread creation/deletion is kernel task = expensive
advantages of threads over processes vanishes

2 possible solutions:

1. scheduler activation, upcall to achieve process switch
2. light-weight processes (LWP)
    user level thread package

Figure 5.1: light-weight processes can run threads

execute scheduler and run thread of parent
may block on systemcall, then other LWP may run
triggered from userspace

Advantages of LWP and user-level thread package:

1. creation, deletion etc is easy, no kernel intervention
2. blocking syscall does not suspend process if enough LWPs are available
3. applications do not see LWP. They only see user-level threads
4. LWP can run on different processors in multiprocessor systems

Disadvantages:

1. LWP creation as expensive as creation of kernel-level thread

Advantages:

- a blocking systemcall blocks only thread, not process $\Rightarrow$ system call for communication in distributed systems

Multiple threads in clients and servers

Clients:

- multiple thread may hide communication delay (distribution transparency)
- web browser opens several connections to load parts of a document/page
- web server may be replicated in same or different location
  $\Rightarrow$ truly parallel access to items and parallel download

Servers:

- single threaded, e.g. file server
  thread serves incoming request, waits for disk, returns file
  serves next
- multithreaded
  dispatcher thread recieves request
  hands over to worker thread
  waits for disk etc.
  dispatcher takes next request
- finite state machine
  only one thread
  examines request, either read from . . . or from disk
  during wait stores requests in table
  serves next request
  manage control either new request or reply from disk
  act accordingly
  process acts as finite state machine that receives messages and acts/changes state

Summary:

model	characteristics
single thread	no parallelism, blocking syscalls
multi thread	parallelism, blocking syscalls
finite state machine	parallelism, non-blocking syscalls

5.1 Virtualisation

V pretends there are more resources then available.

Figure 5.2: virtualisation

Reasons for the need for Virtualization
-hardware changes much faster then SW
$\Rightarrow$ improves portability
-networks consist of different hardware
$\Rightarrow$ enables portability of programs for all
usage (distributed applications, network protocols)

2 Types of Architectures for Virtualisation:

1. Runtime system providing instruction set
   Virtualization of processes (e.g. Wine)

   - interpreted as Java
   - emulated as for Windows applications on UNIX-platform processes VM
2. Virtualisation shields hardware and offers instruction set of the same or other hardware
   - can host different OS that run simultaneosly
   $\Rightarrow$ VMM such as VMware, Xen

5.2 Client-/Serverprocesses

Clients:

Figure 5.3: app specific communication

- b) allows to store data at the server
- **thin client** e.g. X-windows
- thin client should separate application logic from user interaction

Figure 5.4: machine only communication

- often not implemented $\Rightarrow$ poor performance
- compression of interaction commands as solution
- compound documents where user interaction triggers several processing steps on the server. must be implemented (e.g. rotation of picture changes placement in texts)

Servers:

- serves requests on behalf of the client (can server one request at a time)
- Types of servers

  - **iterative Server** handles requests itself
  - **concurrent server** passes requests to worker, e.g. multithreaded server
- server listens to port, endpoint to the client; some ports are reserved for special services

Figure 5.5: listener server

- superserver listens to several ports, replacinf several (mostly idle) servers

Figure 5.6: superserver

16

- stateless servers, keeps no information on state of client → change state without informing the client, e.g. web server

Figure 5.7: stateless server

- soft state server, maintains client state for limited time, e.g. servers informing about updates
- stateful server keeps information about client (file server keeps (client, file) table), often better performance, fault-tolerance poorer
- cookies allow to share information for server upon next visit client sends it'S cookies, allows state information for stateless server

Distributed Servers

Figure 5.8: distributed server

- servers in different locations that have different ip-adresses in DNS under the same name
- MIPv6: mobility support for IPv6
- mobile node has home network with stable home adress (HoA)
- special router is home agent and takes care of traffic to the mobile node
- mobile node receives care-of-adress (CoA), never seen by client
- route optimisation avoids routing through home agent

5.3   Code Migration

- Code migration on (running) process - Why?
- service placement in distributed system ⇒ minimize communication cost

17

- load balancing in multiprocessor machine or cluster $\Rightarrow$ performace
- (security)

Models of Migration

- or process model
    1. code segment, instructions
    2. resource segement, references to external resources, ie.e. file, printer, devices
    3. execution segement, execution state process, stack, private data, programm counter
- **Migration types**
    - weak mobility, transfer code, (1), mabe 3)), which executes from beginning (i.e. java applets)
    - strong mobility, transfer 1)3), stop executions, transfer, resume

Migrating resource segment 2) is difficult

Consider process to resource binding

1. binding by identifier, URL, ftp-server-name
2. binding by value, libraries for programming
3. binding by type, local device, monitor

Resource-machine-binding

1. unattached
2. fastend
3. fixed

pass tp resource binding	unattached	fastened	fixed
by identifier	MV	GR(or MV)	GR
by value	CP	GR(or CP)	GR
by type	RB	RB(or GR,CP)	RB(or GR)

MV:move, GR, global reference, CP: copy value, RB: rebind to locally available resource

Chapter 6

Communication

- Communication in distributed systems is always based on low-level message passing as offered by the underlying network
- message passing is harder than using primitives based on shared memory, as in nondistributed systems
- low-level communication facilities of computer networks are in many ways not suitable due to their lack of distribution transparency.

6.1   RPC - Remote Procedure Call

- allow programs to call procedures located on other machines
- When a process on machine A calls' a procedure on machine B, the calling process on A is suspended, and execution of the called procedure takes place on B.
- Remote procedure call uses stubs to pack parameters in message
- client stub: packs the parameters into a message and requests that message to be sent to the server
- server stub: transforms requests coming in over the network into local procedure calls
- No message passing at all is visible to the programmer
- neither client nor server need to be aware of the intermediate steps or the existence of the network

A remote procedure call occurs in the following steps:

1. The client procedure calls the client stub in the normal way.

2. The client stub builds a message and calls the local operating system.

3. The client's as sends the message to the remote as.

4. The remote as gives the message to the server stub.

5. The server stub unpacks the parameters and calls the server.

6. The server does the work and returns the result to the stub.

7. The server stub packs it in a message and calls its local as.

8. The server's as sends the message to the client's as.

9. The client's as gives the message to the client stub.

10. The stub unpacks the result and returns to the client.

Parameter Marshaling

parameter marshaling: packing parameters into a message is called

Passing Value Parameters

- values are packed into messages (client) and unpacked from messages (server)
- transfered byte-by-byte
- as long as the client and server machines are identical this model works fine
- in a large distributed system, it is common that multiple machine types are present
- $\Rightarrow$ problems because of different character encoding (EBCDIC vs ASCII), represetation of integers (one's complement vs two's complement) or endianness (little endian vs. big endian)

Passing Reference Parameters

- extremly difficult
- pointers are meaningful only within the address space of the process in which it is being used
- replace with copy/restore: copy the datastructure, send it to the server, work on it, send it back, restore at the client

6.2 Asynchronous RPC

Figure 6.1: a: synchronous b: asynchronous RPC

- in conventional procedure calls, when a client calls a remote procedure, the client will block until a reply is returned
- asynchronous RPCs: the server immediately sends a reply back to the client the moment the RPC request is received. Reply acts as an acknowledgment.
- client will continue without further blocking as soon as it has received the server's acknowledgment

- Examples: transferring money from one account to another, adding entries into a database, starting remote services, batch processing...
- Asynchronous RPCs can also be useful when a reply will be returned but the client doesn't need to wait for it and can do nothing in the meantime
- One-Way RPCs: the client does not wait for an acknowledgment from the server
- deferred synchronous RPC: organize the communication between the client and server through two asynchronous RPCs
- foo

6.3 Message oriented communication

General Idea: avoid synchronous communication which blocks sender (RPC)

6.3.1 Message-Oriented Transient Communication

transient: flüchtig, vorrübergehend

Berkeley Sockets

A socket is a communication end point to which an application can write data that are to be sent out over the underlying network, and from which incoming data can be read. A socket forms an abstraction over the actual communication end point that is used by the local operating system for a specific transport protocol.

Figure 6.2: Connection-oriented communication pattern using sockets

- socket: create a new communication end point
- bind: attach a local addres to a socket
- listen: announce willingness to accept connections
- accept: block caller until a connection request arrives
- connect: actively attemt to establish a connection
- send: send some data over the connection
- receive: receive some data over the connection
- close: release the connection

Message-passing-interface (MPI)

- standad for message passing

- designed for parallel applications
- communication within groups of processes
- A $(groupID, processID)$ pair uniquely identifies the source or destination of a message (used instead of a transport-level address)

6.3.2 Message-Oriented Persistent Communication

aka Message-queuing-system, Message-oriented-middleware (MoM)

Figure 4·20. The general organization of a message-queuing system with routers.

Figure 6.3: general organization of a message-queuing system with routers

- asynchronous persistent communication
- offer intermediate-term storage capacity for messages, without requiring either the sender or receiver to be active during message transmission
- transfer may take minutes, not milliseconds
- applications communicate by inserting messages into queues
- messages are only put into and read from local queues
- the message-queuing system takes care that messages are transferred from their source to their destination queue
- message carries destination address
- queue managers
  - a queue manager interacts directly with the application that is sending or receiving a message

22

      – also special queue managers that operate as routers, or relays: they forward incoming messages to other queue managers
- message brokers transform type A into type B, using a set of rules
  - application-level gateway in a message-queuing system
  - convert incoming messages so that they can be understood by the destination application
  - transform messages of type A into type B, using a set of rules
- Examples: Email, workflow, batch processing, queries accross several databases

6.4   stream oriented communication

- temporal relationship between items important
- multimedia data is compressed
- QoS is important

  - bit rate
  - max delay for session setup
  - max end-to-end delay
  - max delay variance (jitter)
  - max round trip delay

- networking solution such as differentiated services
- synchronisation of streams

6.5   Multicast communication

- application level multicast uses overlay
- tree, unique path between each pair of nodes
- mesh, more robust, fault-tolerant

Example: Construct overlay tree for chord
- node that wants to start multicast generates key 128bit/160bit (nid) randomly
- lookup of succ(nid) finds node responsible for key mid
  $\Rightarrow$ succ(nid) becomes root of tree
- join: lookup (nid) creates lookup message with join request routed from P to succ(nid)
- request is forwarded Q (first time forward), Q becomes forwarder
  $\Rightarrow$ P child of Q
- request is first time forwarded by R, R becomes forwarder
  $\Rightarrow$ Q becomes child of R
- multicast: lookup(nid) sends message to the root
  multicast from root

Efficiency?

Quality of application level tree

1. Link stress, number of traversals of same link per packet
2. stretch, relative delay penalty (RDP)
   $\frac{\text{transmission time in overlay}}{\text{transmission time in delay/network}} \Rightarrow$ minimize aggregated stretch, average RDP over all note pairs
3. tree cost, minimize aggregated link cost, link cost = cost between end points
   $\Rightarrow$ find minimal spanning tree

Figure 6.4: example of a overlay network

6.6 Gossip-based-communication

- epidemic behaviour
- a node does not have new data (susceptible), it has the data (infected) or is unwilling to spread (removed)
  Anti-entropy-model
  P chooses randomly Q

  1. P pushes its data to Q
  2. P pulls Q's data
  3. P and Q exchange data

- if many nodes are infected probabiltry for selecting susceptible node is low
  $\Rightarrow$ low probability of data dissemination

- pull works when many nodes are infected. Susceptible node determines spread. They have a high probability to contact infected nodes
- if only one node is infected push/pull is best
- Round is period in which each node at least once selects a neighbor
  number of rounds needed to spread $\approx \mathcal{O}(\log(N))$, N is number of nodes

Rumor spreading, gossiping:
function of nodes that never obtain data: $s = e^{-(k+1)(1-s)}$
e.g. $k = 4, ln(s) = 4, 97$

$\Rightarrow s = 0,007$

less than $0,7$ remain without data

removing data is difficult : delete message is send via gossiping

Chapter 7

Naming

7.1 Flat naming

7.1.1 Distributed Hash Tables

- m-bit identifier (128 or 160 Bit)
- entity with key k is under jurisdiction of node with smallest identifier id $\geq$ k
  $\Rightarrow$ succ(k)
- resolve key k to address of succ(k)
- option 1: each node p keeps succ(p), pred(p) node forwards request for key k to a neighbor
  if pred(p) < k $\leq$ p, return(p)
  $\Rightarrow$ not scalable
- better solution: each Chord node maintains <u>finger table</u> of lenght m

  $$\forall\, 1 \leq i \leq m : FT[i] = succ(p + 2^{i+1}) \mod 2^m$$

  FT[i]=succ($p + 2^{i-1}$) = succ($p + 1$) = succ(2) (smallest id, sucht that id $\geq$ 2)
  i-th entry points to 2^{i-1} ahead of p
- to lookup k node p forwards request to p with index j in ps finger table:
  $q = FT_p[j] \leq k \leq FT_p[j+1]$
- example:
  resolve k = 26 from node 1
  $k = 26 > FT_1[5] \Rightarrow$ forward request to node
  $18 = FT_1[5]$
    - node 18 selects node $20 FT_{18}[2] \leq k < FT_{18}[3]$

    - node 20 selects node $21 \Rightarrow 28$ which is responsible for key 26

    - lookup generally requires $O(log(N))$ steps, N nodes in system
    - join/leave is rather simple
    - keeeping figer table up to date is expensive
- example2:

30	
1	37
2	37
3	37
4	38
5	45
6	0

Figure 7.1: Ring

Superpeers: 2^k, therefore 2^{m-k} normal nodes.

Probability that a node is a supernode: $n \cdot \frac{2^k}{2^m} = 2^{k-m} \cdot n$, with m number of bits, k is number of bits that mark superpeers, n is number of nodes (not maximum possible but the actual number).

Number of superpeers to be expected: $E(x) = n \cdot p$

Chapter 8

Synchronistation

8.1 Clock synchronisation algorithms

System model: each machine has timer that causes H interrupts per second
- clock C adds up ticks (interrupts)
- $C_p(t)$ is clock time on machine p
- perfect clock: $\forall\, p, t : C_p(t) = t$
  $\iff C'_p(t) = \frac{dC_p(t)}{dt} = 1$
  $\widehat{=}$ frequency of clock C_p at time t

Figure 6-5. The relation between clock time and UTe when clocks, tick at different rates.

Figure 8.1: fast, slow & perfect clock

- $C'_p(t) - 1 \widehat{=}$ skew of p's clock, difference to perfect clock.
- $C_p(t) - t \widehat{=}$ offset
- real timers do not interrupt exactly H times per second
- maximum drift ρ is a constant guaranteed/specified by the vendor
  $1 - \rho \leq \frac{dC_p(t)}{dt} \leq 1 + \rho$

- at time Δt after two clocks were synchronized the drift can be max:
  $|C_2(\Delta t) - C_1(\Delta t)| \leq 2\rho\Delta t$
  - if the difference should never exceed δ then synchronisation every $\frac{\delta}{2\rho}$ seconds is needed
- time always moves forward.
  example: given values for C_1 and C_2:

$C_1 = \rho = 0.001$
$C_2 = \rho = 0.001$
$\frac{dC}{dt} = \frac{1}{2*0.001} = \frac{1}{0.002} = \frac{1}{2} * 10^3$
Clocks C_1 and C_2 need to be synchronized every $500s$ to keep time in the same second.

8.2 Network Time Protcol (NTP)

- nodes contact time server that has an accurate clock
- time server passive

Figure 8.2: ntp

A estimates its offset to B as
$\Theta = \frac{(T_2 - T_1) + (T_4 - T_3)}{2}$
assuming communication time is symmetric
delay:
$\delta = (T_4 - T_1) - (T_3 - T_2) = (T_2 - T_1) + (T_4 - T_3)$

- A probes B, B probes A
- NTP stores 8 pairs (Θ, δ) per node pair using $\min(\delta)$ for smallest delay
- either A or B can be more stable
- reference node has <u>stratum 1</u> (clock has stratum 0) (stratum = # Servers to a reference clock)
- lower stratrum level is better, will be used.

8.3 Berkeley algorithm

- assumes no node has 'good' time
- time server polls all nodes for their time
- takes average and adjusts speed of nodes correspondingly
- all nodes agree on time, which may not be correct

8.4 Lamports Logical Clocks

- logical time need not correct in real time.
- needs 'happens before' relation a $\to$ b

Figure 8.3: (a) The time daemon asks all the other machines for their clock values (b) The machines anwer. (c) The time daemon tells everyone how to adjust their clock.

1. if a,b are events in the same process and a happens before b, than a → b is true
2. if a denotes the event of sending a message and b the event of receiving this message by another process then a → b is true

- happens before is transitive:
  $a \rightarrow b \wedge b \rightarrow c \Rightarrow a \rightarrow c$
- concurrency:
  if x, y happen in different processes and neither $x \rightarrow y$ nor $y \rightarrow x$, then x, y are concurrent (which means, it is not known which one comes first)
- ∀ events a, we can assign it a time $C(a)$ on which all processes aggree.
- if $a \rightarrow b$ then $C(a) < C(b)$

  if $C(a) < C(b)$ then not $a \xrightarrow{?} b$

  if $C(a) \not< C(b)$ then $a \not\rightarrow b$

- 4 properties of logical time
  1. No two events get assigned the same time.
  2. Logical times of events in each process are strictly increasing
  3. logical time of sendevent is strictly smaller than receive event for the same message
  4. for any $t \in T$ only finetely many events get assigned logical times smaller then t.
- Example:

Algorithm

C_i = local counter of process P_i

1. Before executing an event (sending a msg over the network, delivering a msg to an app, some internal events) P_i executes
   $C_i \leftarrow C_i + 1$
2. When process P_i sends message m to P_j it sets the timestamp $ts(m)$ of m to the current time $ts(m) \leftarrow C$.
3. upon receipt of a message m process P_j adjust its time to $C_j \leftarrow \max C_j, ts(m)$, then executes step 1 and delivers message

8.4.1 Totally ordered multicast

Examples
- Consider a bank with two data centers A and B, that need to be kept consistent. Each request uses the nearest copy.
  Assume a customer has $1000,- in his bank account and decides to add $100,- using copy A. At the same time 1% interest is added to copy B. What happens? How can we solve the problem?

30

Figure 6-9. (a) Three processes, each with its own clock. The clocks run at different rates. (b) Lamport's algorithm corrects the clocks.

Figure 8.4: ...

- every message is sent to all receivers+itself with timestamp
- Assumption 1: messages from the same sender are received in the order they were sent, and that no messages are lost
- Assumption 2: no two events get assigned the same time
- When msg is received
    1. put it into a local priority queue ordered to the msgs timestamp
    2. multicast an acknowledgment to all other processes (following Lamport $\Rightarrow ts(msg) < ts(ack)$)
- eventually all queues are identical $\Rightarrow$ total order
- Process delivers a queued msg to the app only when
    1. msg is at head of queue
    2. ack received from every process

8.5 Vector Clocks

- Problem with Lamport: does not capture causality
- By construction, we know that for each message $T_{sent}(m_i) < T_{recv}(m_i)$. But what can we conclude in general from $T_{recv}(m_i) < T_{sent}(m_j)$ [Lamport]
- In the case $m_i = m_1$ and $m_j = m_3$ we know at P_2 that m_j was sent after m_i was received. This *may* indicate that sending of m_j has something to do with the receiving of m_i
- Vector Clocks:
  $VC(a) < VC(b)$ means, that event a is known to causally precede event b.
- Each process P_i maintains a Vector VC_i with the following properties:
    1. $VC_i[i]$ is the number of events that occured so far at P_i
       $VC_i[i]$ is the logial clock at P_i
    2. if $VC_i[j] = k$ then P_i knows, that k events have occured at P_j. It is thus P_i's knowledge of the local time at P_j
- To maintain properties:
    1. Before executing an event (i.e., sending a message over the network, delivering a message to an application, or some other internal event), P_i executes $VC_i[i] = VC_i[i] + 1$

Figure 6-12. Concurrent message transmission using logical clocks.

2. When process P_i sends a message m to P_j, it sets m's (vector) timestamp $ts(m)$ equal to VC_i after having executed the previous step
3. Upon the receipt of a message m, process P_j adjusts its own vector by setting $VC_j[k] = max(VC_j[k], ts(m)[k])$ for each k, after which it executes the first step and delivers the message to the application.

- timestamp $ts(m)$ tells the receiver how many events in other processes have preceded the sending of m, and on which m may causally depend.

8.5.1 Casually-ordered multicast

- Casually-ordered mulsticast is weaker than totally ordered multicast
- delivery of message m from P_i to P_j to application layer will be delayed until:
    1. $ts(m)[i] = VC_j[i] + 1$
       = m is the next message that P_j was expecting from process P_i
    2. $ts(m)[k] \leq VC_j[k]\ \ \forall\ k \neq i$
       = P_j has seen all the messages that have been seen by P_i when it sent message m
- Better to be implemented on application layer, because the app knows which messages are causally related.

Figure 6-13. Enforcing causal communication.

Chapter 9

Mutual Exclusion

Access to shared resources
2 types of algorithms: token-based and permission-based
- token is simple, reliability problem (lost token)
- permission difficult in distributed systems

9.1 Centralized algorithm
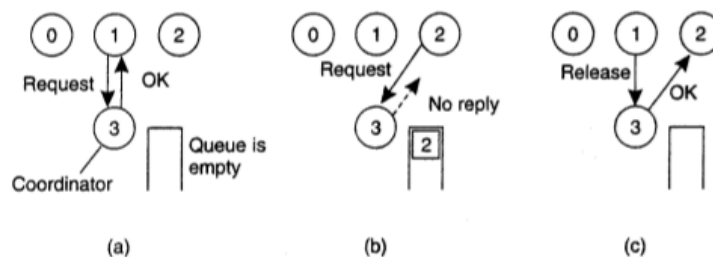
Figure 6-14. (a) Process 1 asks the coordinator for permission to access a shared resource. Permission is granted. (b) Process 2 then asks permission to access the same resource. The coordinator does not reply. (c) When process 1 releases the resource, it tells the coordinator, which then replies to 2.

- one process is coordinator
- coordinator alloes access onl to one process
- fair, requests are processed in order of arrival
- no starvation
- easy to implement
- coordinator is single point of failure
- (handle message loss with ack)
- dead coordinator looks like permission denied

9.2 Decentralized algorithm

- Let m be the count of the coordinators, that answered and allowd entry to the critical section

- Each resource is replicated n times , rname_i is the name of the replica

- each replica has it's own controller, the name is a hash of the rname_i
- if rname is known, each process can generate the address of the controllers
- access to resource when $m > n/2$ controllers grant it
- Let p probability that a coordinator resets during Δt
  $P[k] = prob\{k$ out of m coordinators reset during $\Delta t\} = \left(\binom{m}{k}\right) p^k (1-p)^{m-k}$
- at least $2m - n \geq n + 2 - n = 2$ coordinators need to reset in order to violate the voting. This
  happens with probability $\sum\limits_{k=2m-n}^{n} P[k]$
  e.g. $\Delta t = 10s, n = 32, m = 0,75n$
  Probability of violation in 10^{-40}
- if a process gets less than m votes access to the resource is denied
- random backoff, retry
  many requests, noone gets access
- heavy load $\Rightarrow$ drop in utilisation

9.3 A distributed algorithm

- deterministic
- uses total ordering of events
- process that wants to access a resource sends out message containing (resourcename, process no, current localtime) to all other processes and itself
- process receives a message. Either:

  1. returns OK, if does not want a resource
  2. queues request, if it has resource
  3. compares timestamps, sends OK if timestamp is smallest, queues request and sends no reply else

- grants mutual exclusion without deadlocks or starvation

Figure 6-15. (a) Two processes want to access a shared resource at the same moment., (b) Process 0 has the lowest timestamp. so it wins. (c) When process 0 is done, it sends an *OK* also, so 2 can now go ahead.

Problems:
- node failure $\Rightarrow$ dito
- load, all processes take part in decisions (needs 2(n-1) messages for n processes)
- algorithm is slower, more complicated, more expensive, less robust than centralised alg.
- not a good algorithm

9.4   Token Ring Algorithm

- processes form a logical ring
- token circulates
- owner of token can access resource
- simple and efficient
- not fair under heavy load

Problems:

- token loss -> regenerating?
- crashing nodes

9.5   Comparison

Algorithm	messages per entry/exit	Delay before access	Problems
Centralised	3	2	coordinator crash
Decentralised	$3mk$	$2m$	starvation, low efficiency
Distributed	$2(n-1)$	$2(n-1)$	crash of any process
Token Ring	1 to ∞	0 to ∞	lost token, process crash, fairness?

Chapter 10

Leader Election algorithms

10.1 leader election in a synchronnous ring

Network is a graph G consisting of n nodes connected by unidirectional links. Use $\mod n$ for labels

- elected node is "leader"
- leader election is not possible for identical processes/nodes
- $\rightarrow$ processes have unique id (UID)

10.1.1 LCR algorithm

(Le Lann, Chang, Roberts)

- unidirectional communication
- ring size unknown
- only leader produces output
- algorithm compares UID

```
1 For each node
2   a = a UID, initially i's UID
3   send = a UID or NULL, initially i's UID
4   status = {unknown, leader} initially unknown
5
6 message generation
7   send = current value of send to node i+1
8
9 state transitions
10   send = NULL
11   if incoming message is v (a UID) then
12     v>u:  send v
13     v=u:  status=leader
14     v<u:  do nothing
```

Correctness

Let max index of process with $max(UID)$ let u_{max} is its UID

Figure 10.1: LCR algorithm

Show:

   (i) process max outputs "leader" after n rounds
  (ii) no other process does the same
        We clarify:

 (iii) After n rounds status$_{max}$=leader
and

 (iv) For $0 \leq r \leq n-1$ after r rounds
        $$send_{max} = u_{max}$$
find UID at distance r from i_{max} as it has t og once around.

   Show (iv) for all r: Induction
then (iii)

Complexity

   • time complexity id n rounds
   • communocation complexity $\mathcal{O}(n^2)$
   • not very expensive in time many messages

10.1.2   Algorithm of Hirschberg and Sinclair (HS-Alg)

- reduces number of messages to $\mathcal{O}(n \log n)$

```
1 each process has states with components
2   u, UID: initially i's UID
3   send+ containing NULL or (UID, flag{in, out}, hopcount): initially
      (i's UID, out, 1)
4   send− as send+
```

```
5    status $\in$E{unknown, leader} initailly unknown phase $\in \mathbb
        {N}$: initially 0

6

7  message generation
8    send current send+ to process i+1
9    send current send− to process i−1

10

11  state transitions
12    send+=NULL
13    send−=NULL
14    if message from (i−1) is (v, out, h) then
15      v>u $\land$ h>1: send+ = (v,out,h−1)
16      v>u $\land$ h=1: send− = (v,in,1)
17      v=u status = leader
18    if message from i+1 is (v, out, h) then
19      v>u $\land$ h>1: send− = (v,out, h−1)
20      v>u $\land$ h=1: send+ = (v,in,1)
21      v=u status=leader
22    if message from i−1 is (v,in,1) $\land v\neq u$ then
23      send+=(v,in,1)
24    if message from i+1 is (v,in,1) $\land v\neq u$ then
25      send−=(v,in,1)
26    if both messages from i−1 and i+1 are (u,in,1) then
27      phase++
28      send+=(u,out, $2^{phase}$)
29      send−=(u,out, $2^{phase}$)
```

Complexity

Total number of phases is at most $1 + \lceil \log(n) \rceil$

the total number of messages is at mostin $(1 + \lceil (\log(n)) \rceil \approx \mathcal{O}(n \log n)$

Total time complexity is at most $3n$ if n power of 2 other wise is $5n$

10.1.3   Time slice algorithm

- ring size n is known
- unidirectional
- elects minimum

```
1  phases with n rounds
2  in phase r consisting of rounds (v−1)n+1,\dots,vn
3  only a token carrying UID v is permitted
4  if a process  with UID v exists, then it elects itself leader and
        sends a token wit it's␣UID
```

Complexity: number of messages is n, time complexity $n \cdot u_{min}$

10.1.4   Variable speeds algorithm

- each process i creates a token to tracel around the ring, carrying UID u of origin
- tokens travel at diffeneed speed

Figure 10.2: Timeslice algorithm

- token carrying UID v travels 1 messages every 2^v rounds
- each process memorises smallestUID
- return to origin elects UID

Complexity

. . .

How many messages in total? $\sum\limits_{k=1}^{n} \frac{1}{2^{k-1}} (< 2n)$

Time complexity: $n \cdot 2^{u_{min}}$

10.2   Leader election in a wireless environment

- consider time needed for communication
- every node can initiate an election
- the select is based on information like battery lifetime or processing power
- the nodes only participate in one election, also there are more than one initiating nodes

Process to choose a leader:

Look at figure 10.3

1. one node starts the election and sends ELECT to its neighbours
2. the node, which sended the ELECT message becomes the parent of the node
3. the node sends the request to all neighbours except the parent node
4. the node acknowledges the parent after all the children acknowledged the election

5. if a node recieves an ELECT message, but has already a parent, it acknowledges immediately
6. if a the node has only neighbours, which are the parent or have other parents, the node becomes a leaf and sends back its value
7. the waiting nodes send the node id and the highest value of the childs to the parent, after all acknowledged
8. in the end, the initiating node knows the node with the highest value and broadcast to all nodes, that this is the leader

10.3   The Bully Algorithm(flooding) (Garcia-Mdina, 1982)

- process P holds election
1. P sends ELECT message to all processes
2. P wins if there is no response $\Rightarrow$ P is leader and sends COORDINATOR message to all available nodes.
3. if Q answers with OK, Q takes over and sends ELECT again.

Figure 6-22. Election algorithm in a wireless network, with node *a* as the source. (a) Initial network., (b)-(e) The build-tree phase (last broadcast step by nodes f and i not shown). (f) Reporting of best node to source.

Figure 10.3: Election in wireless networks

Figure 6-20. The bully election algorithm. (a) Process 4 holds an election. (b) Processes 5 and 6 respond. telling 4 to stop. (c) Now 5 and 6 each hold an election. (d) Process 6 tells 5 to stop. (e) Process 6 wins and tells everyone.

Figure 10.4: The bully election algorithm

Chapter 11

Consistency and Consensus

Anm.: Quelle ist hier der deutsche Tanenbaum/Steen, daher sind manche englischen Begriffe evtl. nicht korrekt gewählt.

- Distributed Systems use replication of data to improve *performance* and/or *reliability*
- replication for scalability
  How to keep replicas consistent?
  Many types of consistency

  - data-centric-constistency
  - client-centric-consistency
  - monotic reads: successive reads return the same or newer value
  - monitic write: a write op must be completed before the next write by the same process
  - read-your-own-write: write is always seen by read of same process
  - write-follows-read: write on previous read takes place on the same or more recent value
- Do not discuss replica placement
- Object-replication

11.1 Data-centric consistency

The more strict a model is, the stronger are the assumptions, that may be drawn from a series of events, but the harder to implement.

11.1.1 Strict consistency

Every read of x returns the value of the last write on x

Example strict consistency

P1: W(x)a

P2:   R(x)a

Non strict consistent example

P1: W(x)a

P2:   R(x)Null R(x)a

Simple and clear idea, but this definition requires existence of global time (hard to realise), so that the *last* write on a variable is unambiguous. This requirement is very hard to accomplish.

If a data storage fulfills this requirement, then every write must be visible globally (for every process in the system).

11.1.2   Sequential/linear consistency

Weak version of strict consistency, irrelevance of the chronological order of write events. That means, that there is no global time or chronology, but some order of the events is globally visible and the same for every process. Linear consistency just adds the requirement, that the events are ordered by a globally (mayhaps ambiguous) timestamp.

Example sequential consistency

P1: W(x)a

P2:   W(x)b

P3:   R(x)b   R(x)a

P4:   R(x)b   R(x)a

Example non-sequential consistency

P1: W(x)a

P2:   W(x)b

P3:   R(x)b   R(x)a

P4:   R(x)a   R(x)b

11.1.3   Causal consistency

Weak version of sequential consistency, only causally dependent events must appear strictly ordered on a global level. Two causally dependent events may be a write event *W(x)a*, a read event *R(x)a* followed by a write *W(x)b*, then the write *W(x)b* may be based on the value writte in *W(x)a*.

Example causal consistency

P1:	W(x)a			W(x)c	
P2:		R(x)a	W(x)b		
P3:		R(x)a			R(x)c	R(x)b
P4:		R(x)a			R(x)b	R(x)c

Example non-causal consistency

P1:	W(x)a				
P2:		R(x)a	W(x)b		
P3:				R(x)b	R(x)a
P4:				R(x)a	R(x)b

11.1.4 FIFO-consistency

Weak version of causal consitency. The writes of a single process (which are ordered and causally dependent) appear in this exact order globally, but the sequence of writes between multiple processes is not strictly ordered. In other words, only the casual dependencies inside one process appear in the correct order, dependencies between different processes may be mixed.

Example FIFO-consistency

P1:	W(x)a					
P2:		R(x)a	W(x)b	W(x)c		
P3:					R(x)b	R(x)a	R(x)c
P4:					R(x)a	R(x)b	R(x)c

11.1.5 Weak consistency

The idea of weak consistency is, that also the relatively weak assertions of FIFO-consistency often are too hard to implement or just unnecessary. The globally visible actions are reduced to critical sections (or areas) where the actions become globally visible only after leaving the critical section, so only the results of the actions become visible. To realise that behaviour a mechanism to synchronise on the current state of a synchronisation variable S is introduced, such that processes who are interested in the current status of this variable may update their data, but other processes are not bothered with updates.

In this consistency scheme, consistency is not defined absolutely but only temporarily (at synchronization points).

Example weak consistency

P1: W(x)a	W(x)b	S			

P2:			R(x)a	R(x)b	S

P3:			R(x)b	R(x)a	S

Example non-weak consistency

P1: W(x)a	W(x)b	S	

P2:		S	R(x)a

11.1.6 Release-consistency

In the weak-consistency-scheme there is no difference between synchronising after a write (i.e. a write-through to all replications) and synchronising before reading data, but these to scenarios require different actions on lower layers. Two additional procedures are added *acquire(..)* and *release(..)*.

11.1.7 Entry-consistency

Similar to release-consistency, but synchronisation variables (and procedures acq(..), rel(..)) are responsible for a designated set of variables, not for all variables in the global scope.

11.2 Client-centric consistency

11.2.1 Eventual consistency

11.3 Reliable multicast protocols

- atomic multicast requirement
  all requests arrive at all servers in the same order

11.3.1 Distributed Commit

- an operatio is performed by group or non of the nodes of the group
- reliable multicast operation = delivery of message
- distributed transaction: operation = execution of transaction
- uses coordinator
- one-phased commit

- two-phase commit (2PC) (Jim Gray, 1978)
  - distributed transaction involves several processors each on a different machine
    2 phases with each 2 steps:

    1. coordinator $\xrightarrow{vote\ request}$ all participants

2. participant $\xrightarrow[vote\ abort]{vote\ commit}$ coordinator
3. if all commit
   coordinator $\xrightarrow{global-commit}$ all participants
   else
   coordinator $\xrightarrow{global-abort}$ all participants

4. if commit, then participants locally commit
   else participants locally abort

Problems if failures occur

   * coordinator blocks in: wait
   * participant blocks in: ready, init
   $\Rightarrow$ blocking commit protocol
 - use timeouts to unblock
 - repeat request
 - in state ready P con contact Q
   if Q is in contact, then coordinator died after sending to Q
   before sending to P $\Rightarrow$ P can commit
   if Q is in abort $\Rightarrow$ abort
   if Q is in init $\Rightarrow$ abort
   if Q is in ready $\rightarrow$ abort or no decision contact R

- Three-phase commit (3PC) (Steen, 1981)

  - avoids blocking in the presence of fail-stop crashes
  - states satisfy the following conditions

    1. there is no state from which directly follow commit or abort follows
    2. there is no state in which it is not possible to make a final decision and from which
       a transaction to a commit state can be made
    $\Rightarrow$ necessary and sufficient conditions for non-blocking commit protocol
  - abort branch as in 2PC
  - blocking states: paritcipant: init -> abort
    coordinator: wait -> abort
    precommit, knowing P voted for commit
    $\Rightarrow$ global-commit+recovery of P
    participant: ready
    coordinator failed as in 2PC
    precommit: contact other participants: if Q in precommit $\Rightarrow$ commit
    if Q is in init $\Rightarrow$ abort
  - Q can be in INIT only if no participant is in precommit
  - participant can reach precommit only if coordinator was in precommit already
  - In 2PC a crashed participant could recover to commit, while all others are still in ready
  - if one process is in ready recovery can be only to states ready, init, abort, precommit
    $\Rightarrow$ surviving processes can come to final solution
- Paxos (Leslie Lamport, late 80s)
  - does not block with at most n/2-1 failures
    Paxos adds to 2PC:

* ordering of proposals
* majority voting for acceptance

Duelling proposer

1. Aufgabe: Terminiologie (wichtige Konzepte, erklären, vergleichen, bla,bla) dann durch die Themen des Semsters, übungszettel, gerne ausrechnen (Fingertable, Metriken von overlaynetzen, komplexität von protokollen (wie viele nachrichten braucht ein protokoll), logische uhren (stellen oder so)), erklären, Peersimaufgabe(n) (programm angucken, was macht das programm?, überblick, wie modifizieren für fkt x, cycle driven vs event driven (was wofür)), last auf dem netz, kein gnuplot programm auf papier!!! Hilfsmittel: mitbringen, was man will, außer internet, telefon, freunde usw...