

Lecture Notes
Distributed System

Hinnerk van Bruinehsen
Tobias Höppner
Tobias Famulla
Lord Pimmelgurke
Los Jufflos
JensTM

SoSe 2013

Contents

1	Verteilte Systeme/Distributed Systems	3
1.1	Orga	3
1.1.1	Elektisches	3
1.1.2	Übungen	3
1.1.3	Material/Inhalt	3
2	Distributed Systems	5
2.1	Transparency	5
2.2	Open System	6
2.3	Scalability	6
2.4	Problems	6
2.4.1	Scalability Limitations	6
2.4.2	Geographical Scalability	6
2.4.3	Security Problems	6
2.5	Scaling techniques	6
2.6	pitfalls	6
2.7	Types of distributed systems	7
3	Architectures of distributed Systems	8
3.1	System architectures	9
4	PeerSim	12
5	Processes	13
5.1	Virtualisation	15
5.2	Client-/Serverprocesses	16
5.3	Code Migration	17
6	Communication	19
6.1	RPC - Remote Procedure Call	19
6.2	Asynchronous RPC	20
6.3	Message oriented communication	21
6.3.1	Message-Oriented Transient Communication	21
6.3.2	Message-Oriented Persistent Communication	22
6.4	Stream Oriented Communication	23
6.5	Multicast communication	24
6.5.1	Application Level Multicasting	24
6.5.2	Gossip-based-communication	25

7	Naming	27
7.1	Flat naming	27
7.1.1	Distributed Hash Tables	27
8	Synchronisation	29
8.1	Clock synchronisation algorithms	29
8.2	Network Time Protocol (NTP)	30
8.3	Berkeley algorithm	30
8.4	Lamports Logical Clocks	30
8.4.1	Totally ordered multicast	32
8.5	Vector Clocks	32
8.5.1	Casually-ordered multicast	33
9	Mutual Exclusion	34
9.1	Centralized algorithm	34
9.2	Decentralized algorithm	34
9.3	A distributed algorithm	35
9.4	Token Ring Algorithm	36
9.5	Comparison	36
10	Leader Election algorithms	37
10.1	leader election in a synchronous ring	37
10.1.1	LCR algorithm	37
10.1.2	Algorithm of Hirschberg and Sinclair (HS-Alg)	38
10.1.3	Time slice algorithm	42
10.1.4	Variable speeds algorithm	42
10.2	Leader election in a wireless environment	43
10.3	The Bully Algorithm(flooding) (Garcia-Mdina, 1982)	44
11	Consistency and Consensus	46
11.1	Data-centric consistency	46
11.1.1	Strict consistency	46
11.1.2	Sequential/linear consistency	47
11.1.3	Causal consistency	47
11.1.4	FIFO-consistency	48
11.1.5	Weak consistency	48
11.1.6	Release-consistency	49
11.1.7	Entry-consistency	49
11.2	Client-centric consistency	49
11.2.1	Eventual consistency	49
11.2.2	Monotone read-consistency	49
11.2.3	Monotone write-consistency	49
11.2.4	Read your writes	50
11.2.5	Writes follow reads	50
11.2.6	Implementing client-side consistency	50
11.3	Reliable multicast protocols	50
11.3.1	Distributed Commit	50

Chapter 1

Verteilte Systeme/Distributed Systems

1.1 Orga

VL Di 10-12 (nicht am 23.04.)
Ue Do 10-12

1.1.1 Elektisches

- (kvv)
- Website AG
- Sakai

1.1.2 Übungen

- ca. 5 Übungsblätter, 14-tägig
- Vorträge in Gruppen über „verteilte Systeme“

1.1.3 Material/Inhalt

1. Hälfte Distributed Systems (Tanenbaum, van Steen)
 - Architektur
 - Prozesse
 - Kommunikation
 - Namen
 - Synchronisation
 - Konsistenz
 - Replikation
 - Fehlertoleranz
2. Hälfte Distributed Algorithms (Nancy Lynch)
 - synchronous network algorithms
 - network models (leader election, shortest path, distributed consensus, byzantine agreement)

- asynchronous network algorithms (shared memory, mutual exclusion, resource allocation, consensus)
- timing
- network resource allocation
- failure detectors

Chapter 2

Distributed Systems

Def: A distributed System is a collection of independent computers that appears to it's users as a single coherent system.

Characteristics:

- autonomous components
- appears as single system
- communication is hidden
- organisation is hidden
(could be high-performance mainframe or sensor net)
- heterogenous system offers homogenous look/interface

Objectives:

- provide resources (printer, storage, computing)
 - share in a controlled, efficient way
 - grant access \Rightarrow connect users and resources

2.1 Transparency

hide the fact that processes and resources are physically distributed.

Types of transparency:

access hide differences in data representation and how a resource is accessed

location hide where a resource is located

migration hide that a resource may move to another location

relocation hide that a resource may be moved to another location while in use

replication hide that a resource is replicated

concurrency hide that a resource may be shared by serveral competitive users

failure hide the failure and recovery of a resource

transparency is desirable, but not always perfectly possible

tradeoff between transparency and complexity, maintainability and performance

2.2 Open System

- service interfaces specified using Interface Definition Language (IDL)
- service specification as text

2.3 Scalability

is an important property, distributed systems should be scalable in

size number of nodes, users, resources

geographic spread

administration

2.4 Problems

2.4.1 Scalability Limitations

centralized services A single server for all users

centralized data A single on-line telephone book

centralized algorithms Doing routing based on complete information

2.4.2 Geographical Scalability

design existing distributed systems were designed for LANs

communication LAN-based systems often use synchronous communication and is inherently unreliable (virtually always p2p)

2.4.3 Security Problems

selfprotection from malicious attacks from the new domain

domainprotection the domain protects itself from attacks from a new distrib. sys.

2.5 Scaling techniques

hiding communication latencies use only asynchronous communication

distribution split components into smaller parts :)

replication of components, caching, enables load balancing

2.6 pitfalls

1. reliable network
2. secure network
3. homogenous network
4. constant topology
5. zero latency
6. infinite bandwidth
7. zero transport cost
8. one administrator!

2.7 Types of distributed systems

- computing systems
 - cluster computing

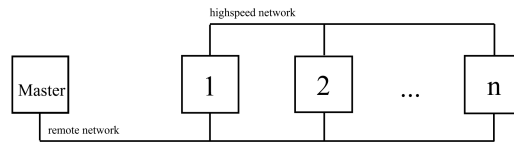


Figure 2.1: cluster computing

- grid computing(virtual organisation, geographically distributed and heterogenous))
- distributed information systems
 - transaction processing systems (database)
 - ACID** (atomicity, consistency, isolation, durability)
 - Atomic To the outside world, the transaction, happens indivisibly
 - Consistent The transaction does not violate system invariants
 - Isolated Concurrent transactions do not interfere with each other
 - Durable Once a transaction commits, the changes are permanent
 - enterprise systems
- Distributed pervasive systems
 - small, wireless, adhoc, no administration
 - home automation, health systems, sensor networks

Why do we need distributed systems?

- performance
- distribution inherent
- reliability
- incremental growth (scalability)
- sharing resources

Chapter 3

Architectures of distributed Systems

- how to split software into components
⇒ Softwarearchitecture
 - how to build a system out of the components
⇒ Systemarchitecture
- Middleware can help to create distribution transparency

Architecturestyles:

- Layered architecture
⇒ network stack, messages or data flow up and down
 - control flow between layers
 - requests down
 - reply up
- Object-based architectures
 - interaction between components
 - e.g. remote procedure calls
 - can be client-server system
- data-centered architectures
 - data is key element
 - communication over data, distributed database
 - web-systems mostly data-centric
- event-based architecture
 - publish-subscribe systems
 - processes communicates through events
 - publisher announces events at broker
 - ⇒ loose coupling (publisher and subscriber need not to know each other), decoupled in space
 - ⇒ scalability better than client-server, parallel processing, caching

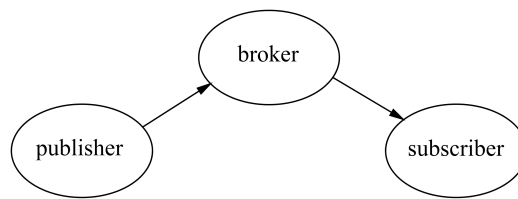


Figure 3.1: publish subscribe system

Event-based and data-based can be combined
 ⇒ shared Data space

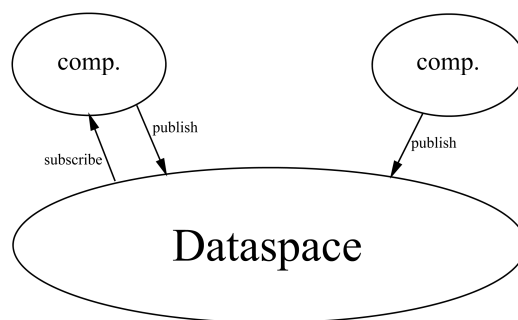


Figure 3.2: shared data space

3.1 System architectures

1. centralized architectures
 - client - server

- (i) single point of failure
- (ii) performance (server is bottleneck)

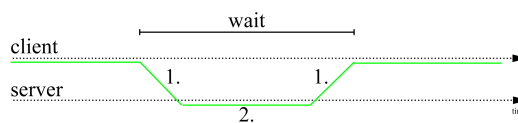


Figure 3.3: client server simple waiting situation

- (a) communication problems
- (b) server problems

can request be repeated without harm?

⇒ request is idempotent

(iii) application layering

Layers:

- 1.) User interface
- 2.) processing
- 3.) data level




Figure 3.4: application layer

⇒ a lot of waiting

⇒ does not scale

2. Decentralized architectures

- vertical distribution (layering)
different logic on different machines
- horizontal distribution
replicated client/server operating on different data
⇒ overlay-underlay hides physical structure by adding logical structure

Structured P2P architectures

- most popular technique is distributed hash tables (DHT)
 - randomly 128 bit or 160 bit ke for data and nodes. Two or more keys are very unlikely
 - Chord system arranges items in a ring
 - data item k is assigned to node with smallest identifier $id \geq k$
- ie item 1 belongs to node 1
item 2 belongs to node 2
for each item k_i $succ(k)=id$
returns the name of the node k is assigned to
to find data item k the function $LOOKUP(k)$ returns the adress of $succ(k)$ in $O(\log(N))$ (later!)

membership management

join:

create SHA1 identifier

$LOOKUP(id) = succ(id)$

contact $succ(id)$ and $pred(id)$ to join ring

leave:

node id informs $succ(id)$ and $pred(id)$ and assigns it's data to $succ(id)$

Content adressable network (CAN)

- d-dimensional cartesian space
- every node draws random number
- space is divided among nodes
- every data draws identifier (coordinates) which assigns a node
- join
 - select random point
 - half the square in which id falls
 - assign item to centers
- leave
 - one node takes the rectangle

⇒ reassign rectangles periodically

Unstructured P2P Network

- random graph
 - each node maintains a list of c neighbours
 - partial view or neighbourhood list with age
 - nodes exchange neighbour information
- active thread
select peer

PUSH

select $c/2$ youngest entries+myself
send to peer

PULL

receive peer buffer
construct new partial view
increment age

passive thread
recieve buffer from peer

PULL:

select $c/2$
send to peer
construct new partial view increment age

Chapter 4

PeerSim

Chapter 5

Processes

processes

- execution of program
- processor creates virtual processor
- for each program everything is stored in process table
- transparent sharing of resources,(processor, memory) separation
- each virtual processor has it's own independent address space
- process switch is expensive, (save cpu context, pointers, translation lookaside buffer (TLB), memory management unit (MMU))
- perhaps even swaps to disk, if memory exhausted

Summary:

- Several threads share one process.
- Processes have own address space, threads don't.
- Processes and threads each have userland and kernel implementations

2 possible solutions:

1. scheduler activation, upcall to achieve process switch
2. light-weight processes (LWP)
user level thread package

threads

- several threads share CPU
- thread context has little memory information, perhaps mutex lock
- threads avoid blocking application (e.g. spreadsheet, computation of dependent cells, intermediate backup)
- thread switch is fast
- user level threads allow parallel computation of program sections
- I/O or other blocking system calls block all threads, but thread creation/deletion is kernel task = expensive
- advantages of threads over processes vanishes




Figure 5.1: light-weight processes can run threads

execute scheduler and run thread of parent
may block on syscall, then other LWP may run
triggered from userspace

Advantages of LWP and user-level thread package:

1. creation, deletion etc is easy, no kernel intervention
2. blocking syscall does not suspend process if enough LWPs are available
3. applications do not see LWP. They only see user-level threads
4. LWP can run on different processors in multiprocessor systems

Disadvantages:

1. LWP creation as expensive as creation of kernel-level thread

Advantages:

- a blocking syscall blocks only thread, not process \Rightarrow system call for communication in distributed systems

Multiple threads in clients and servers

Clients:

- multiple thread may hide communication delay (distribution transparency)
- web browser opens several connections to load parts of a document/page
- web server may be replicated in same or different location
 \Rightarrow truly parallel access to items and parallel download

Servers:

- single threaded, e.g. file server
thread serves incoming request, waits for disk, returns file
serves next
- multithreaded
dispatcher thread receives request
hands over to worker thread
waits for disk etc.
dispatcher takes next request
- finite state machine
only one thread
examines request, either read from . . . or from disk
during wait stores requests in table
serves next request
manage control either new request or reply from disk
act accordingly
process acts as finite state machine that receives messages and acts/changes state

Summary:

model	characteristics
single thread	no parallelism, blocking syscalls
multi thread	parallelism, blocking syscalls
finite state machine	parallelism, non-blocking syscalls

5.1 Virtualisation




Figure 5.2: virtualisation

V pretends there are more resources than available.

Reasons for the need for Virtualization

- hardware changes much faster than SW
- ⇒ improves portability
- networks consist of different hardware
- ⇒ enables portability of programs for all usage (distributed applications, network protocols)

2 Types of Architectures for Virtualisation:

1. Runtime system providing syscall API/ABI
Virtualization of processes (e.g. Wine)
 - syscalls are interpreted (like Java)
 - emulated as for Windows applications on UNIX-platform processes VM
2. Virtualisation shields hardware and offers instruction set of the same or other hardware
 - can host different OS that run simultaneously
 - ⇒ VMM such as VMware, Xen

5.2 Client-/Serverprocesses

Clients:




Figure 5.3: a) app specific communication




Figure 5.4: b) machine only communication

- b) allows to store data at the server
- **thin client** e.g. X-windows
- thin client should separate application logic from user interaction
- often not implemented \Rightarrow poor performance
- compression of interaction commands as solution
- compound documents where user interaction triggers several processing steps on the server. must be implemented (e.g. rotation of picture changes placement in texts)

Servers:

- serves requests on behalf of the client (can server one request at a time)
- Types of servers
 - **iterative Server** handles requests itself
 - **concurrent server** passes requests to worker, e.g. multithreaded server
- server listens to port, endpoint to the client; some ports are reserved for special services
- superserver listens to several ports, replacing several (mostly idle) servers
- stateless servers, keeps no information on state of client \rightarrow change state without informing the client, e.g. web server
- soft state server, maintains client state for limited time, e.g. servers informing about updates




Figure 5.5: listener server




Figure 5.6: superserver

- stateful server keeps information about client (file server keeps (client, file) table), often better performance, fault-tolerance poorer
- cookies allow to share information for server upon next visit client sends it'S cookies, allows state information for stateless server

Distributed Servers

- servers in different locations that have different ip-adresses in DNS under the same name
- MIPv6: mobility support for IPv6
- mobile node has home network with stable home adress (HoA)
- special router is home agent and takes care of traffic to the mobile node
- mobile node receives care-of-adress (CoA), never seen by client
- route optimisation avoids routing through home agent

5.3 Code Migration

- Code migration on (running) process - Why?
- service placement in distributed system \Rightarrow minimize communication cost
- load balancing in multiprocessor machine or cluster \Rightarrow performance
- (security)

Models of Migration

- or process model
 1. code segment, instructions
 2. resource segment, references to external resources, i.e. file, printer, devices
 3. execution segment, execution state process, stack, private data, programm counter
- **Migration types**
 - weak mobility, transfer code, (1), mabe 3)), which executes from beginning (i.e. java applets)
 - strong mobility, transfer 1)3), stop executions, transfer, resume




Figure 5.7: stateless server




Figure 5.8: distributed server

Migrating resource segment 2) is difficult
Consider process to resource binding

1. binding by identifier, URL, ftp-server-name
2. binding by value, libraries for programming
3. binding by type, local device, monitor

Resource-machine-binding

1. unattached
2. fastend
3. fixed

pass tp resource binding	unattached	fastened	fixed
by identifier	MV	GR(or MV)	GR
by value	CP	GR(or CP)	GR
by type	RB	RB(or GR,CP)	RB(or GR)

MV:move, GR, global ref-

erence, CP: copy value, RB: rebind to locally available resource

Chapter 6

Communication

- Communication in distributed systems is always based on low-level message passing as offered by the underlying network
- message passing is harder than using primitives based on shared memory, as in nondistributed systems
- low-level communication facilities of computer networks are in many ways not suitable due to their lack of distribution transparency.

6.1 RPC - Remote Procedure Call

- allow programs to call procedures located on other machines
- When a process on machine A calls a procedure on machine B, the calling process on A is suspended, and execution of the called procedure takes place on B.
- Remote procedure call uses stubs to pack parameters in message
- client stub: packs the parameters into a message and requests that message to be sent to the server
- server stub: transforms requests coming in over the network into local procedure calls
- No message passing at all is visible to the programmer
- neither client nor server need to be aware of the intermediate steps or the existence of the network

A remote procedure call occurs in the following steps:

1. The client procedure calls the client stub in the normal way.
2. The client stub builds a message and calls the local operating system.
3. The client's as sends the message to the remote as.
4. The remote as gives the message to the server stub.
5. The server stub unpacks the parameters and calls the server.
6. The server does the work and returns the result to the stub.
7. The server stub packs it in a message and calls its local as.

8. The server's as sends the message to the client's as.
9. The client's as gives the message to the client stub.
10. The stub unpacks the result and returns to the client.

Parameter Marshaling

parameter marshaling: packing parameters into a message is called

Passing Value Parameters

- values are packed into messages (client) and unpacked from messages (server)
- transferred byte-by-byte
- as long as the client and server machines are identical this model works fine
- in a large distributed system, it is common that multiple machine types are present
- ⇒ problems because of different character encoding (EBCDIC vs ASCII), representation of integers (one's complement vs two's complement) or endianness (little endian vs. big endian)

Passing Reference Parameters

- extremely difficult
- pointers are meaningful only within the address space of the process in which it is being used
- replace with copy/restore: copy the datastructure, send it to the server, work on it, send it back, restore at the client

6.2 Asynchronous RPC

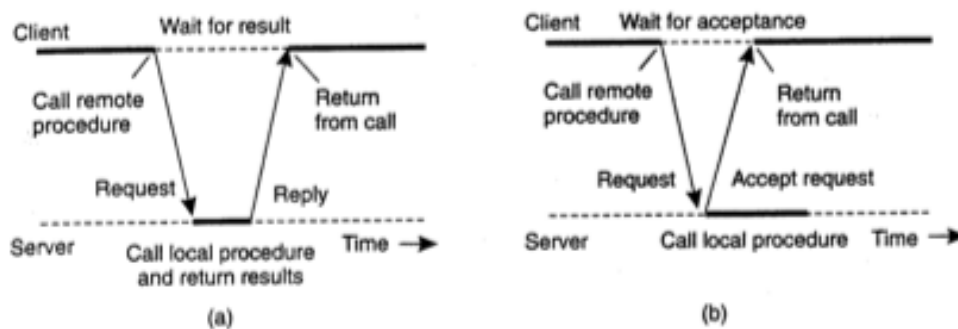


Figure 6.1: a: synchronous b: asynchronous RPC

- in conventional procedure calls, when a client calls a remote procedure, the client will block until a reply is returned
- asynchronous RPCs: the server immediately sends a reply back to the client the moment the RPC request is received. Reply acts as an acknowledgment.
- client will continue without further blocking as soon as it has received the server's acknowledgment

- Examples: transferring money from one account to another, adding entries into a database, starting remote services, batch processing...
- Asynchronous RPCs can also be useful when a reply will be returned but the client doesn't need to wait for it and can do nothing in the meantime
- One-Way RPCs: the client does not wait for an acknowledgment from the server
- deferred synchronous RPC: organize the communication between the client and server through two asynchronous RPCs
- foo

6.3 Message oriented communication

General Idea: avoid synchronous communication which blocks sender (RPC)

6.3.1 Message-Oriented Transient Communication

transient: flüchtig, vorübergehend

Berkeley Sockets

A socket is a communication end point to which an application can write data that are to be sent out over the underlying network, and from which incoming data can be read. A socket forms an abstraction over the actual communication end point that is used by the local operating system for a specific transport protocol.

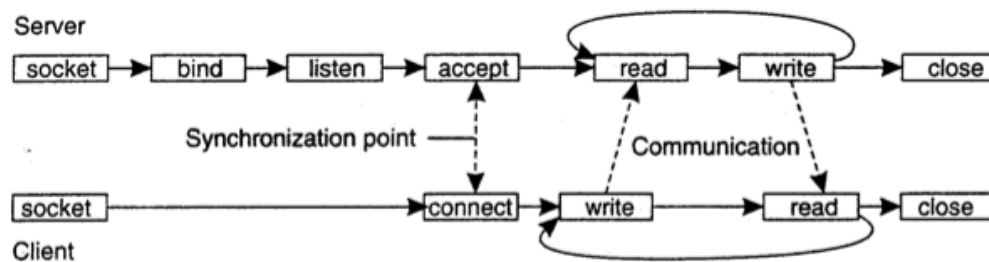


Figure 6.2: Connection-oriented communication pattern using sockets

- socket: create a new communication end point
- bind: attach a local address to a socket
- listen: announce willingness to accept connections
- accept: block caller until a connection request arrives
- connect: actively attempt to establish a connection
- send: send some data over the connection
- receive: receive some data over the connection
- close: release the connection

Message-passing-interface (MPI)

- standard for message passing

- designed for parallel applications
- communication within groups of processes
- A (*groupID, processID*) pair uniquely identifies the source or destination of a message (used instead of a transport-level address)

6.3.2 Message-Oriented Persistent Communication

aka Message-queuing-system, Message-oriented-middleware (MoM)

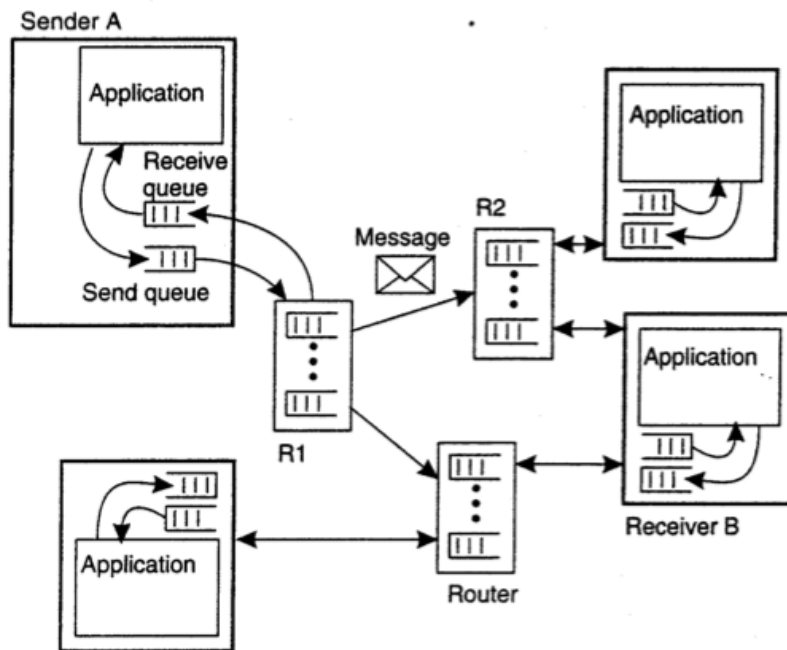


Figure 4-20. The general organization of a message-queuing system with routers.

Figure 6.3: general organization of a message-queuing system with routers

- asynchronous persistent communication
- offer intermediate-term storage capacity for messages, without requiring either the sender or receiver to be active during message transmission
- transfer may take minutes, not milliseconds
- applications communicate by inserting messages into queues
- messages are only put into and read from local queues
- the message-queuing system takes care that messages are transferred from their source to their destination queue
- message carries destination address
- queue managers
 - a queue manager interacts directly with the application that is sending or receiving a message

- also special queue managers that operate as routers, or relays: they forward incoming messages to other queue managers
- message brokers transform type A into type B, using a set of rules
 - application-level gateway in a message-queuing system
 - convert incoming messages so that they can be understood by the destination application
 - transform messages of type A into type B, using a set of rules
- Examples: Email, workflow, batch processing, queries accross several databases

6.4 Stream Oriented Communication

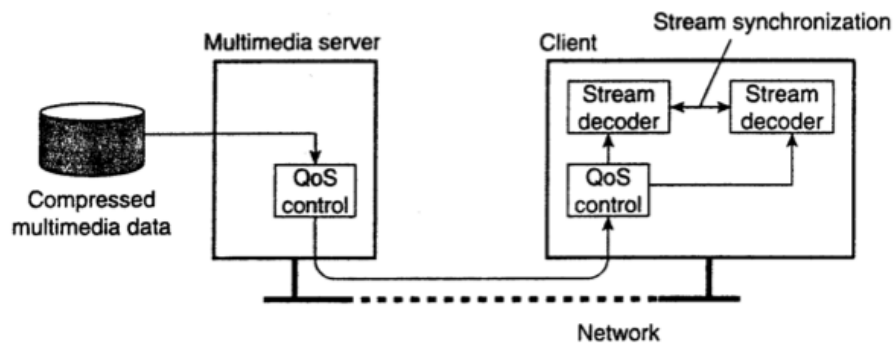


Figure 6.4: A general architecture for streaming stored multimedia data over a network

- form of communication in which timing plays a crucial role
- in continuous media, the temporal relationships between different data items are fundamental to correctly interpreting what the data actually means
- multimedia data will need to be compressed substantially in order to reduce the required network capacity
- simple stream: consists of only a single sequence of data
- complex stream: consists of several related (often time dependent) simple streams (substreams)
- QoS
 1. The required bit rate at which data should be transported.
 2. The maximum delay until a session has been set up (i.e., when an application can start sending data).
 3. The maximum end-to-end delay (i.e., how long it will take until a data unit makes it to a recipient).
 4. The maximum delay variance, or jitter.
 5. The maximum round-trip delay.
- synchronisation of streams
 - Synchronization of streams deals with maintaining temporal relations between streams
 - Synchronization takes place at the level of the data units of which a stream is made up
 - In other words, we can synchronize two streams only between data units
 - Example: Playing a movie in which the video stream needs to be synchronized with the audio

6.5 Multicast communication

6.5.1 Application Level Multicasting

- sending data to multiple receivers
- For many years, this topic has belonged to the domain of network protocol
- With the advent of peer-to-peer technology various application-level multicasting techniques have been introduced
- The basic idea in application-level multicasting is that nodes organize into an overlay network, which is then used to disseminate information to its members
- connections between nodes in the overlay network may cross several physical links, and as such, routing messages within the overlay may not be optimal in comparison to what could have been achieved by network-level routing

Approaches to building the overlay

- tree: nodes may organize themselves directly into a tree, meaning that there is a unique (overlay) path between every pair of nodes
- mesh network: nodes organize into a mesh network in which every node will have multiple neighbors and, in general, there exist multiple paths between every pair of nodes
- mesh network generally provides higher robustness

Example: Construct overlay tree for chord

- node that wants to start multicast generates key 128bit/160bit (mid) randomly
- lookup of $\text{succ}(\text{mid})$ finds node responsible for key mid
⇒ $\text{succ}(\text{mid})$ becomes root of tree
- join multicast: lookup (mid) creates lookup message with join request routed from P to $\text{succ}(\text{mid})$
- request is forwarded by Q (first time forward), Q becomes forwarder
⇒ P child of Q
- request is then forwarded by R, R becomes forwarder
⇒ Q becomes child of R
- if Q or R is already forwarder: no forward
⇒ Q becomes child of R
- multicast: lookup(mid) sends message to the root
multicast from root

Efficiency

- building a tree is not difficult once we have organized the nodes into an overlay
- building an efficient tree may be difficult
- The quality of an application-level multicast tree is generally measured by three different metrics
 1. Link stress: defined per link and counts how often a packet crosses the same link
 2. Stretch / Relative Delay Penalty (RDP)
ratio in the delay between two nodes in the overlay, and the delay that those two nodes




Figure 6.5: The relation between links in an overlay and actual network-level routes

would experience in the underlying network: $\frac{\text{transmission time in overlay}}{\text{transmission time in delay/network}}$
 \Rightarrow minimize aggregated stretch, average RDP over all node pairs

3. tree cost: global metric, generally related to minimizing the aggregated link costs
 link cost = cost between end points
 \Rightarrow find minimal spanning tree

6.5.2 Gossip-based-communication

- epidemic behaviour: model information dissemination in a (large) distributed system after the spreading of infectious diseases
- infected node: a node that holds data that it is willing to spread
- susceptible: a node does not yet have the new data
- removed: an updated node that is not willing or able to spread its data
- we assume we can distinguish old from new data, for example, because it has been timestamped or versioned

Anti-entropy-model

A node P picks another node Q at random, and subsequently exchanges updates with Q. There are three approaches to exchanging updates:

1. P only pushes its own updates to Q
 - updates can be propagated only by infected nodes
 - if many nodes are infected, the probability of each one selecting a susceptible node is relatively small
 - \Rightarrow chances are that a particular node remains susceptible for a long period simply because it is not selected by an infected node
2. P only pulls in new updates from Q
 - spreading updates is essentially triggered by susceptible nodes

- high probability to contact an infected node and pull in the updates
 - \Rightarrow pull-based approach works much better when many nodes are infected
3. P and Q send updates to each other (i.e., a push-pull approach)
- if only one node is infected push/pull is best
-
- Round is period in which each node at least once selects a neighbor
 - number of rounds needed to spread $\approx \mathcal{O}(\log(N))$, N is number of nodes

Rumor spreading, gossiping

- if node P has just been updated for data item x , it contacts an arbitrary other node Q and tries to push the update to Q
- if Q was already updated by another node, P may lose (with some probability) interest in spreading the update any further
- P then becomes removed
-
- Fraction of nodes that never obtain data: $s = e^{-(k+1)(1-s)}$
e.g. $k = 4, \ln(s) = -4,97$
 $\Rightarrow s = 0,007$
less than 0,7% remain without data

Removing Data

- problem: deletion of a data item destroys all information on that item
 \Rightarrow when a data item is simply removed from a node, that node will eventually receive old copies of the data item and interpret those as updates on something it did not have before
- \Rightarrow record the deletion of a data item as just another update, and keep a record of that deletion

Chapter 7

Naming

7.1 Flat naming

7.1.1 Distributed Hash Tables

- m-bit identifier (128 or 160 Bit)
- entity with key k is under jurisdiction of node with smallest identifier $id \geq k$
 $\Rightarrow succ(k)$
- resolve key k to address of $succ(k)$
- option 1: each node p keeps $succ(p)$, $pred(p)$ node forwards request for key k to a neighbor if $pred(p) < k \leq p$, return(p)
 \Rightarrow not scalable
- better solution: each Chord node maintains finger table of length m

$$\forall 1 \leq i \leq m : FT[i] = succ(p + 2^{i-1}) \mod 2^m$$

$FT[i] = succ(p + 2^{i-1}) = succ(p + 1) = succ(2)$ (smallest id, such that $id \geq 2$)
 i -th entry points to 2^{i-1} ahead of p

- to lookup k node p forwards request to p with index j in p 's finger table:
 $q = FT_p[j] \leq k \leq FT_p[j+1]$
- example:
resolve $k = 26$ from node 1
 $k = 26 > FT_1[5] \Rightarrow$ forward request to node
 $18 = FT_1[5]$
 - node 18 selects node 20 $FT_{18}[2] \leq k < FT_{18}[3]$
 - node 20 selects node 21 $\Rightarrow 28$ which is responsible for key 26
 - lookup generally requires $O(\log(N))$ steps, N nodes in system
 - join/leave is rather simple
 - keeping finger table up to date is expensive
- example2:

30	
1	37
2	37
3	37
4	38
5	45
6	0




Figure 7.1: Ring

Superpeers: 2^k , therefore 2^{m-k} normal nodes.
Probability that a node is a supernode: $n \cdot \frac{2^k}{2^m} = 2^{k-m} \cdot n$, with m number of bits, k is number of bits that mark superpeers, n is number of nodes (not maximum possible but the actual number).
Number of superpeers to be expected: $E(x) = n \cdot p$

Chapter 8

Synchronisation

8.1 Clock synchronisation algorithms

System model: each machine has timer that causes H interrupts per second

- clock C adds up ticks (interrupts)
- $C_p(t)$ is clock time on machine p
- perfect clock: $\forall p, t : C_p(t) = t$
 $\iff C'_p(t) = \frac{dC_p(t)}{dt} = 1$
 $\hat{=}$ frequency of clock C_p at time t




Figure 6-5. The relation between clock time and UTC when clocks tick at different rates.

Figure 8.1: fast, slow & perfect clock

- $C'_p(t) - 1 \hat{=}$ skew of p 's clock, difference to perfect clock.
- $C_p(t) - t \hat{=}$ offset
- real timers do not interrupt exactly H times per second
- maximum drift ρ is a constant guaranteed/specified by the vendor
 $1 - \rho \leq \frac{dC_p(t)}{dt} \leq 1 + \rho$

- at time Δt after two clocks were synchronized the drift can be max:
 $|C_2(\Delta t) - C_1(\Delta t)| \leq 2\rho\Delta t$
- if the difference should never exceed δ then synchronisation every $\frac{\delta}{2\rho}$ seconds is needed
- time always moves forward.

example: given values for C_1 and C_2 :

$$C_1 = \rho = 0.001$$

$$C_2 = \rho = 0.001$$

$$\frac{dC}{dt} = \frac{1}{2 \cdot 0.001} = \frac{1}{0.002} = \frac{1}{2} \cdot 10^3$$

Clocks C_1 and C_2 need to be synchronized every 500s to keep time in the same second.

8.2 Network Time Protocol (NTP)

- nodes contact time server that has an accurate clock
- time server passive




Figure 8.2: ntp

A estimates its offset to B as

$$\Theta = \frac{(T_2 - T_1) + (T_4 - T_3)}{2}$$

assuming communication time is symmetric

delay:

$$\delta = (T_4 - T_1) - (T_3 - T_2) = (T_2 - T_1) + (T_4 - T_3)$$

- A probes B, B probes A
- NTP stores 8 pairs (Θ, δ) per node pair using $\min(\delta)$ for smallest delay
- either A or B can be more stable
- reference node has stratum 1 (clock has stratum 0) (stratum = # Servers to a reference clock)
- lower stratum level is better, will be used.

8.3 Berkeley algorithm

- assumes no node has 'good' time
- time server polls all nodes for their time
- takes average and adjusts speed of nodes correspondingly
- all nodes agree on time, which may not be correct

8.4 Lamports Logical Clocks

- logical time need not correct in real time.
- needs 'happens before' relation $a \rightarrow b$
 1. if a,b are events in the same process and a happens before b, then $a \rightarrow b$ is true




Figure 8.3: (a) The time daemon asks all the other machines for their clock values (b) The machines answer. (c) The time daemon tells everyone how to adjust their clock.

2. if a denotes the event of sending a message and b the event of receiving this message by another process then $a \rightarrow b$ is true
- happens before is transitive:
 $a \rightarrow b \wedge b \rightarrow c \Rightarrow a \rightarrow c$
- concurrency:
 if x, y happen in different processes and neither $x \rightarrow y$ nor $y \rightarrow x$, then x, y are concurrent (which means, it is not known which one comes first)
- \forall events a , we can assign it a time $C(a)$ on which all processes agree.
- if $a \rightarrow b$ then $C(a) < C(b)$
 if $C(a) < C(b)$ then not $a \stackrel{?}{\rightarrow} b$
 if $C(a) \not< C(b)$ then $a \not\rightarrow b$
- 4 properties of logical time
 1. No two events get assigned the same time.
 2. Logical times of events in each process are strictly increasing
 3. logical time of sendevent is strictly smaller than receive event for the same message
 4. for any $t \in T$ only finetely many events get assigned logical times smaller then t .
- Example:




Figure 6-9. (a) Three processes, each with its own clock. The clocks run at different rates. (b) Lamport's algorithm corrects the clocks.

Figure 8.4: ...

Algorithm

C_i = local counter of process P_i

1. Before executing an event (sending a msg over the network, delivering a msg to an app, some internal events) P_i executes
 $C_i \leftarrow C_i + 1$
2. When process P_i sends message m to P_j it sets the timestamp $ts(m)$ of m to the current time
 $ts(m) \leftarrow C_i$.
3. upon receipt of a message m process P_j adjust its time to $C_j \leftarrow \max(C_j, ts(m))$, then executes step 1 and delivers message

8.4.1 Totally ordered multicast

Examples

- Consider a bank with two data centers A and B, that need to be kept consistent. Each request uses the nearest copy.
Assume a customer has \$1000,- in his bank account and decides to add \$100,- using copy A. At the same time 1% interest is added to copy B. What happens? How can we solve the problem?
 - every message is sent to all receivers+itself with timestamp
 - Assumption 1: messages from the same sender are received in the order they were sent, and that no messages are lost
 - Assumption 2: no two events get assigned the same time
 - When msg is received
 1. put it into a local priority queue ordered to the msg's timestamp
 2. multicast an acknowledgment to all other processes (following Lamport $\Rightarrow ts(msg) < ts(ack)$)
 - eventually all queues are identical \Rightarrow total order
 - Process delivers a queued msg to the app only when
 1. msg is at head of queue
 2. ack received from every process

8.5 Vector Clocks

- Problem with Lamport: does not capture causality
- By construction, we know that for each message $T_{sent}(m_i) < T_{recv}(m_i)$. But what can we conclude in general from $T_{recv}(m_i) < T_{sent}(m_j)$ [Lamport]
- In the case $m_i = m_1$ and $m_j = m_3$ we know at P_2 that m_j was sent after m_i was received. This may indicate that sending of m_j has something to do with the receiving of m_i
- Vector Clocks:
 $VC(a) < VC(b)$ means, that event a is known to causally precede event b.
- Each process P_i maintains a Vector VC_i with the following properties:
 1. $VC_i[i]$ is the number of events that occurred so far at P_i
 $VC_i[i]$ is the logical clock at P_i
 2. if $VC_i[j] = k$ then P_i knows, that k events have occurred at P_j . It is thus P_i 's knowledge of the local time at P_j
- To maintain properties:
 1. Before executing an event (i.e., sending a message over the network, delivering a message to an application, or some other internal event), P_i executes $VC_i[i] = VC_i[i] + 1$
 2. When process P_i sends a message m to P_j , it sets m 's (vector) timestamp $ts(m)$ equal to VC_i after having executed the previous step




Figure 6-12. Concurrent message transmission using logical clocks.

3. Upon the receipt of a message m , process P_j adjusts its own vector by setting $VC_j[k] = \max(VC_j[k], ts(m)[k])$ for each k , after which it executes the first step and delivers the message to the application.
- timestamp $ts(m)$ tells the receiver how many events in other processes have preceded the sending of m , and on which m may causally depend.

8.5.1 Casually-ordered multicast

- Casually-ordered multicast is weaker than totally ordered multicast
- delivery of message m from P_i to P_j to application layer will be delayed until:
 1. $ts(m)[i] = VC_j[i] + 1$
= m is the next message that P_j was expecting from process P_i
 2. $ts(m)[k] \leq VC_j[k] \quad \forall k \neq i$
= P_j has seen all the messages that have been seen by P_i when it sent message m
- Better to be implemented on application layer, because the app knows which messages are causally related.




Figure 6-13. Enforcing causal communication.

Chapter 9

Mutual Exclusion

Access to shared resources

2 types of algorithms: token-based and permission-based

- token is simple, reliability problem (lost token)
- permission difficult in distributed systems

9.1 Centralized algorithm




Figure 6-14. (a) Process 1 asks the coordinator for permission to access a shared resource. Permission is granted. (b) Process 2 then asks permission to access the same resource. The coordinator does not reply. (c) When process 1 releases the resource, it tells the coordinator, which then replies to 2.

- one process is coordinator
- coordinator allows access only to one process
- fair, requests are processed in order of arrival
- no starvation
- easy to implement
- coordinator is single point of failure
- (handle message loss with ack)
- dead coordinator looks like permission denied

9.2 Decentralized algorithm

- Let m be the count of the coordinators, that answered and allowed entry to the critical section

- Each resource is replicated n times, $rname_i$ is the name of the replica
- each replica has its own controller, the name is a hash of the $rname_i$
- if $rname$ is known, each process can generate the address of the controllers
- access to resource when $m > n/2$ controllers grant it
- Let p probability that a coordinator resets during Δt
 $P[k] = \text{prob}\{k \text{ out of } m \text{ coordinators reset during } \Delta t\} = \binom{m}{k} p^k (1-p)^{m-k}$
- at least $2m - n \geq n + 2 - n = 2$ coordinators need to reset in order to violate the voting. This happens with probability $\sum_{k=2m-n}^n P[k]$
e.g. $\Delta t = 10s, n = 32, m = 0,75n$
Probability of violation in 10^{-40}
- if a process gets less than m votes access to the resource is denied
- random backoff, retry
many requests, no one gets access
- heavy load \Rightarrow drop in utilisation

9.3 A distributed algorithm

- deterministic
- uses total ordering of events
- process that wants to access a resource sends out message containing (resourcename, process no, current localtime) to all other processes and itself
- process receives a message. Either:
 1. returns OK, if does not want a resource
 2. queues request, if it has resource
 3. compares timestamps, sends OK if timestamp is smallest, queues request and sends no reply else
- grants mutual exclusion without deadlocks or starvation




Figure 6-15. (a) Two processes want to access a shared resource at the same moment., (b) Process 0 has the lowest timestamp. so it wins. (c) When process 0 is done, it sends an OK also, so 2 can now go ahead.

Problems:

- node failure \Rightarrow dito
- load, all processes take part in decisions (needs $2(n-1)$ messages for n processes)
- algorithm is slower, more complicated, more expensive, less robust than centralised alg.
- not a good algorithm

9.4 Token Ring Algorithm

- processes form a logical ring
- token circulates
- owner of token can access resource
- simple and efficient
- not fair under heavy load

Problems:

- token loss -> regenerating?
- crashing nodes

9.5 Comparison

Algorithm	messages per entry/exit	Delay before access	Problems
Centralised	3	2	coordinator crash
Decentralised	$3mk$	$2m$	starvation, low efficiency
Distributed	$2(n - 1)$	$2(n - 1)$	crash of any process
Token Ring	1 to ∞	0 to ∞	lost token, process crash, fairness?

Chapter 10

Leader Election algorithms

10.1 leader election in a synchronous ring

Network is a graph G consisting of n nodes connected by unidirectional links. Use $\text{mod } n$ for labels

- elected node is "leader"
- leader election is not possible for identical processes/nodes
→ processes have unique id (UID)

10.1.1 LCR algorithm

(Le Lann, Chang, Roberts)

- unidirectional communication
- ring size unknown
- only leader produces output
- algorithm compares UID
- One or more p_i s can take the initiative and start an election, by sending an election message containing their id to p_{i+1}

```
1 For each node
2   u = a UID, initially i's UID
3   send = a UID or NULL, initially i's UID
4   status = {unknown, leader} initially unknown
5
6 message generation
7   send = current value of send to node i+1
8
9 state transitions
10  send = NULL
11  if incoming message is v (a UID) then
12    v > u: send v
13    v = u: status = leader
14    v < u: do nothing
```

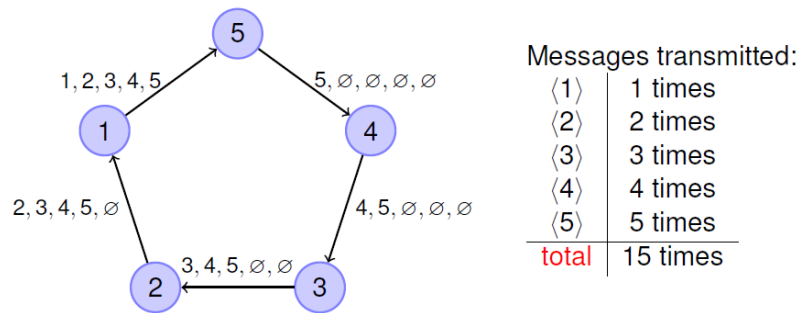


Figure 10.1: LCR algorithm

Correctness

Let $\max UID$ of process with $\max UID$ let u_{\max} is its UID
Show:

- (i) process \max outputs "leader" after n rounds
- (ii) no other process does the same

We clarify:

- (iii) After n rounds $\text{status}_{\max} = \text{leader}$
and

- (iv) For $0 \leq r \leq n - 1$ after r rounds

$$\text{send}_{\max} = u_{\max}$$

find UID at distance r from i_{\max} as it has to go once around.

Show (iv) for all r : Induction
then (iii)

Complexity

- time complexity is n rounds
- communication complexity $\mathcal{O}(n^2)$
- not very expensive in time, but many messages

10.1.2 Algorithm of Hirschberg and Sinclair (HS-Alg)

- reduces number of messages to $\mathcal{O}(n \log n)$

- 1 each process has states with components
- 2 u , UID: initially i 's UID
- 3 send+ containing NULL or (UID, flag{in, out}, hopcount): initially
(i 's UID, out, 1)
- 4 send- as send+
- 5 $\text{status} \in E\{\text{unknown, leader}\}$ initially unknown phase $\in N$: initially 0
- 6
- 7 message generation

```

8   send current send+ to process i+1
9   send current send- to process i-1
10
11 state transitions
12   send+=NULL
13   send-=NULL
14   if message from (i-1) is (v, out, h) then
15       v>u ∧ h>1: send+ = (v, out, h-1)
16       v>u ∧ h=1: send- = (v, in, 1)
17       v=u status = leader
18   if message from i+1 is (v, out, h) then
19       v>u ∧ h>1: send- = (v, out, h-1)
20       v>u ∧ h=1: send+ = (v, in, 1)
21       v=u status=leader
22   if message from i-1 is (v, in, 1) ∧ v ≠ u then
23       send+=(v, in, 1)
24   if message from i+1 is (v, in, 1) ∧ v ≠ u then
25       send-=(v, in, 1)
26   if both messages from i-1 and i+1 are (u, in, 1) then
27       phase++
28       send+=(u, out, 2phase)
29       send-=(u, out, 2phase)

```

Complexity

Total number of phases is at most $8n(1 + \lceil \log(n) \rceil)$

the total number of messages is at most in $(1 + \lceil (\log(n)) \rceil) \approx \mathcal{O}(n \log n)$

Total time complexity is at most $3n$ if n power of 2 otherwise is $5n$




Figure 10.2: HS Phase 1




Figure 10.3: HS Phase 2




Figure 10.4: HS Phase 3

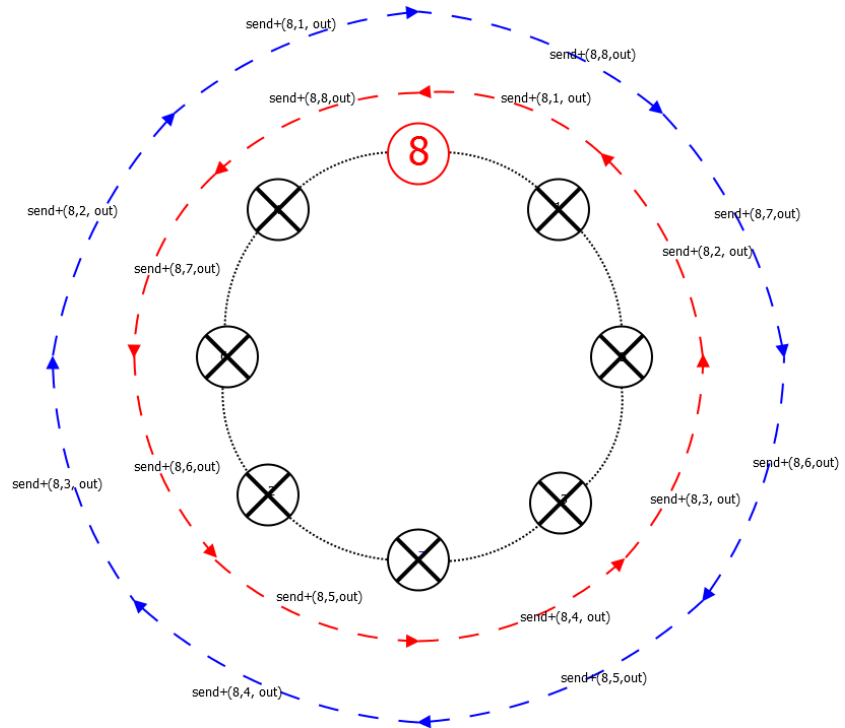


Figure 10.5: HS Phase 4

Show that the number of messages sent by the HS is at most $8n(1 + \lceil \log n \rceil)$, which means it is $O(n \log n)$.

In Phase k können maximal $4 \cdot 2^k$ Nachrichten (in und out) für einen Kandidaten (Knoten, der in seiner Nachbarschaft keinen größeren Knoten hat) existieren.

In Phase $k = 0$ gibt es n Kandidaten $\Rightarrow 2n$ out Nachrichten und maximal $\frac{n}{2} \cdot 2$ in Nachrichten $= 3n$

In Phase $k \geq 1$ gibt es maximal noch $\left\lfloor \frac{n}{2^{k-1}+1} \right\rfloor$ Kandidaten. $2^{k-1} + 1$ ist der minimale Abstand zwischen zwei Gewinnern aus Phase $k - 1$

\Rightarrow die Anzahl der Nachrichten, die in Phase k gesendet werden ist $4 \cdot 2^k \cdot \left\lfloor \frac{n}{2^{k-1}+1} \right\rfloor = 8n \cdot \left\lfloor \frac{2^{k-1}}{2^{k-1}+1} \right\rfloor$

Offensichtlich ist die maximale Anzahl der Phasen ist $\lceil \log n \rceil + 1$

Offensichtlich werden in der letzten Phase $2n$ Nachrichten gesendet

$$3n + \sum_{k=1}^{\lceil \log n \rceil - 1} \left(8n \cdot \left\lfloor \frac{2^{k-1}}{2^{k-1}+1} \right\rfloor \right) + 2n \leq 5n + 8n \cdot (\lceil \log n \rceil - 1)$$

10.1.3 Time slice algorithm

- ring size n is known
- unidirectional
- elects minimum

```

1 phases with n rounds
2 in phase r consisting of rounds  $(v-1) \cdot n + 1, \dots, vn$ 
3 only a token carrying UID  $v$  is permitted
4 if a process with UID  $v$  exists, then it elects itself the leader
   and sends a token with it's UID

```




Figure 10.6: Timeslice algorithm

Complexity: number of messages is n , time complexity $n \cdot u_{min}$

10.1.4 Variable speeds algorithm

- each process i creates a token to travel around the ring, carrying UID u of origin

- tokens travel at different speed
- token carrying UID v travels 1 messages every 2^v rounds
- each process memorises smallest UID
- return to origin elects UID

Complexity

...

How many messages in total? $\sum_{k=1}^n \frac{1}{2^{k-1}} (< 2n)$

Time complexity: $n \cdot 2^{uid_{min}}$

10.2 Leader election in a wireless environment

- consider time needed for communication
- every node can initiate an election
- the select is based on information like battery lifetime or processing power
- the nodes only participate in one election, also there are more than one initiating nodes

Process to choose a leader:

Look at figure 10.7

1. one node starts the election and sends ELECT to its neighbours
2. the node, which send the ELECT message becomes the parent of the node
3. the node sends the request to all neighbours except the parent node
4. the node acknowledges the parent after all the children acknowledged the election
5. if a node receives an ELECT message, but has already a parent, it responds immediately, that another node is its parent
6. if a node has only neighbours, which are the parent or have other parents, the node becomes a leaf and sends back its value
7. the waiting nodes send the node id and the highest value of the children to the parent, after all acknowledged
8. in the end, the initiating node knows the node with the highest value and broadcast to all nodes, that this is the leader




Figure 6-22. Election algorithm in a wireless network, with node *a* as the source. (a) Initial network. (b)-(e) The build-tree phase (last broadcast step by nodes *f* and *i* not shown). (f) Reposting of best node to source.

Figure 10.7: Election in wireless networks

10.3 The Bully Algorithm(flooding) (Garcia-Mdina, 1982)

- process P holds election

1. P sends ELECT message to all processes with higher number
2. P wins if there is no responses \Rightarrow P is leader and sends COORDINATOR message to all available nodes.
3. if Q (with higher number) answers with OK, Q takes over and sends ELECT to all higher nodes again.




Figure 6-20. The bully election algorithm. (a) Process 4 holds an election. (b) Processes 5 and 6 respond, telling 4 to stop. (c) Now 5 and 6 each hold an election. (d) Process 6 tells 5 to stop. (e) Process 6 wins and tells everyone.

Figure 10.8: The bully election algorithm

Chapter 11

Consistency and Consensus

Anm.: Quelle ist hier der deutsche Tanenbaum/Steen, daher sind manche englischen Begriffe evtl. nicht korrekt gewählt.

- Distributed Systems use replication of data to improve *performance* and/or *reliability*
- replication for scalability
 - How to keep replicas consistent?
 - Many types of consistency
 - data-centric-consistency
 - client-centric-consistency
 - monotonic reads: successive reads return the same or newer value
 - monitic write: a write op must be completed before the next write by the same process
 - read-your-own-write: write is always seen by read of same process
 - write-follows-read: write on previous read takes place on the same or more recent value
- Do not discuss replica placement
- Object-replication

11.1 Data-centric consistency

The more strict a model is, the stronger are the assumptions, that may be drawn from a series of events, but the harder to implement.

11.1.1 Strict consistency

Every read of x returns the value of the last write on x

Example strict consistency

P1: $W(x)a$

P2: $R(x)a$

Non strict consistent example

P1: $W(x)a$

P2: $R(x)\text{Null}$ $R(x)a$

Simple and clear idea, but this definition requires existence of global time (hard to realise), so that the *last* write on a variable is unambiguous. This requirement is very hard to accomplish.

If a data storage fulfills this requirement, then every write must be visible globally (for every process in the system).

11.1.2 Sequential/linear consistency

Weak version of strict consistency, irrelevance of the chronological order of write events. That means, that there is no global time or chronology, but some order of the events is globally visible and the same for every process. Linear consistency just adds the requirement, that the events are ordered by a globally (mayhaps ambiguous) timestamp.

Example sequential consistency

P1: $W(x)a$

P2: $W(x)b$

P3: $R(x)b$ $R(x)a$

P4: $R(x)b$ $R(x)a$

Example non-sequential consistency

P1: $W(x)a$

P2: $W(x)b$

P3: $R(x)b$ $R(x)a$

P4: $R(x)a$ $R(x)b$

11.1.3 Causal consistency

Weak version of sequential consistency, only causally dependent events must appear strictly ordered on a global level. Two causally dependent events may be a write event $W(x)a$, a read event $R(x)a$ followed by a write $W(x)b$, then the write $W(x)b$ may be based on the value written in $W(x)a$.

Example causal consistency

P1:	W(x)a	W(x)c		
P2:	R(x)a	W(x)b		
P3:	R(x)a		R(x)c	R(x)b
P4:	R(x)a		R(x)b	R(x)c

Example non-causal consistency

P1:	W(x)a		
P2:	R(x)a	W(x)b	
P3:		R(x)b	R(x)a
P4:		R(x)a	R(x)b

11.1.4 FIFO-consistency

Weak version of causal consistency. The writes of a single process (which are ordered and causally dependent) appear in this exact order globally, but the sequence of writes between multiple processes is not strictly ordered. In other words, only the casual dependencies inside one process appear in the correct order, dependencies between different processes may be mixed.

Example FIFO-consistency

P1:	W(x)a			
P2:	R(x)a	W(x)b	W(x)c	
P3:		R(x)b	R(x)a	R(x)c
P4:		R(x)a	R(x)b	R(x)c

11.1.5 Weak consistency

The idea of weak consistency is, that also the relatively weak assertions of FIFO-consistency often are too hard to implement or just unnecessary. The globally visible actions are reduced to critical sections (or areas) where the actions become globally visible only after leaving the critical section, so only the results of the actions become visible. To realise that behaviour a mechanism to synchronise on the current state of a synchronisation variable S is introduced, such that processes who are interested in the current status of this variable may update their data, but other processes are not bothered with updates.

In this consistency scheme, consistency is not defined absolutely but only temporarily (at synchronization points).

Example weak consistency

P1: W(x)_a W(x)_b S

P2: R(x)_a R(x)_b S

P3: R(x)_b R(x)_a S

Example non-weak consistency

P1: W(x)_a W(x)_b S

P2: S R(x)_a

11.1.6 Release-consistency

In the weak-consistency-scheme there is no difference between synchronising after a write (i.e. a write-through to all replications) and synchronising before reading data, but these two scenarios require different actions on lower layers. Two additional procedures are added *acquire(..)* and *release(..)*.

11.1.7 Entry-consistency

Similar to release-consistency, but synchronisation variables (and procedures *acq(..)*, *rel(..)*) are responsible for a designated set of variables, not for all variables in the global scope.

11.2 Client-centric consistency

In contrast to the data-centric consistency models client-centric consistency models do not try to maintain a system wide consistent view, but rather try to hide inconsistencies on client level.

11.2.1 Eventual consistency

(popular example: Domain Name Service) Usually big, distributed, replicated databases, which tolerate a high amount of inconsistency, but eventually become consistent when there are no updates to the database. It is important that clients access the system via fixed nodes, otherwise client updates may take a while, until they are visible for the client itself which is a clear violation of local consistency (read your writes). Remember that in this section the consistency view is only local/client-specific!

11.2.2 Monotone read-consistency

When a datum x is read by a process, then every following read of x results in the same value, or in newer value.

11.2.3 Monotone write-consistency

When a process writes to a datum x , then all previous write operations have to be finished/applied.

11.2.4 Read your writes

The result of a previous write operation of a process on a datum x is always visible in a following read of the same process

11.2.5 Writes follow reads

A write by a process to a datum x , which follows a previous read of x , is performed on this read value, or a newer one.

11.2.6 Implementing client-side consistency

The idea here is to monitor the client's reads and writes and add an ID to each operation. When reading/writing to a server the current status of the client (a list of the reads/writes) is added to the query and the server checks, whether the local storage of the client has to be updated before executing the query.

11.3 Reliable multicast protocols

- atomic multicast requirement
all requests arrive at all servers in the same order

11.3.1 Distributed Commit

- an operation is performed by group or non of the nodes of the group
- reliable multicast operation = delivery of message
- distributed transaction: operation = execution of transaction
- uses coordinator
- one-phased commit
 - distributed transaction involves several processors each on a different machine
 - 2 phases with each 2 steps:

1. coordinator $\xrightarrow{\text{vote request}}$ all participants
2. participant $\xrightarrow[\text{vote abort}]{\text{vote commit}}$ coordinator
3. if all commit
coordinator $\xrightarrow{\text{global-commit}}$ all participants
else
coordinator $\xrightarrow{\text{global-abort}}$ all participants
4. if commit, then participants locally commit
else participants locally abort




Figure 11.1: 2PC

5. two-phase commit (2PC) (Jim Gray, 1978)

Problems if failures occur

- * coordinator blocks in: wait
- * participant blocks in: ready, init
- ⇒ blocking commit protocol
- use timeouts to unblock
- repeat request
- in state ready P con contact Q
 - if Q is in contact, then coordinator died after sending to Q
 - before sending to P ⇒ P can commit
 - if Q is in abort ⇒ abort
 - if Q is in init ⇒ abort
 - if Q is in ready → abort or no decision contact R

• Three-phase commit (3PC) (Steen, 1981)

- avoids blocking in the presence of fail-stop crashes
- states satisfy the following conditions

1. there is no state from which directly follow commit or abort follows
2. there is no state in which it is not possible to make a final decision and from which a transaction to a commit state can be made

⇒ necessary and sufficient conditions for non-blocking commit protocol

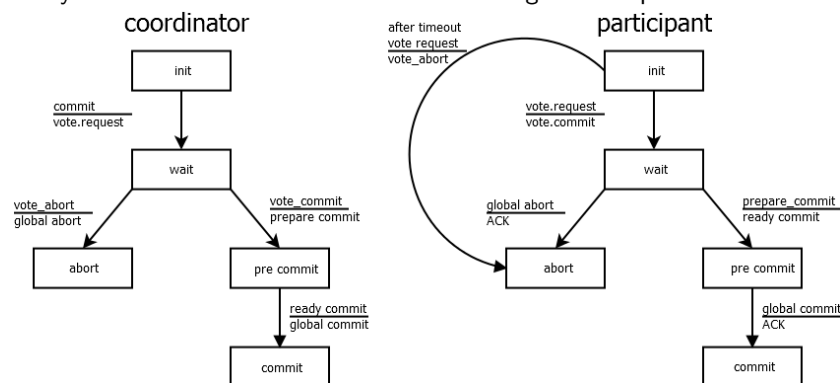


Figure 11.2: 3PC

- abort branch as in 2PC
 - blocking states: participant: init -> abort
 coordinator: wait -> abort
 precommit, knowing P voted for commit
 ⇒ global-commit+recovery of P
 participant: ready
 coordinator failed as in 2PC
 precommit: contact other participants: if Q in precommit ⇒ commit
 if Q is in init ⇒ abort
 - Q can be in INIT only if no participant is in precommit
 - participant can reach precommit only if coordinator was in precommit already
 - In 2PC a crashed participant could recover to commit, while all others are still in ready
 - if one process is in ready recovery can be only to states ready, init, abort, precommit
 ⇒ surviving processes can come to final solution
 - Paxos (Leslie Lamport, late 80s)
 - does not block with at most $n/2-1$ failures
 - Paxos adds to 2PC:
 - * ordering of proposals
 - * majority voting for acceptance
 - Duelling proposer
1. Aufgabe: Terminiologie (wichtige Konzepte, erklären, vergleichen, bla,bla) dann durch die Themen des Semsters, übungszettel, gerne ausrechnen (Fingertable, Metriken von overlaynetzen, komplexität von protokollen (wie viele nachrichten braucht ein protokoll), logische uhren (stellen oder so)), erklären, Peersimaufgabe(n) (programm angucken, was macht das programm?, überblick, wie modifizieren für fkt x, cycle driven vs event driven (was wofür)), last auf dem netz, kein gnuplot programm auf papier!!! Hilfsmittel: mitbringen, was man will, außer internet, telefon, freunde usw...